



# Using PinPlay for Reproducible Analysis and Replay Debugging

**Harish Patil, Cristiano Pereira, Charles(Chuck) Yount**

**Intel Corporation**

**Milind Chabbi, Rice University**

**With contributions from:**

**Mack Stallcup, Greg Lueck, Tevi Devor, David Wootton (Intel  
Corporation)**

# Tutorial Objective

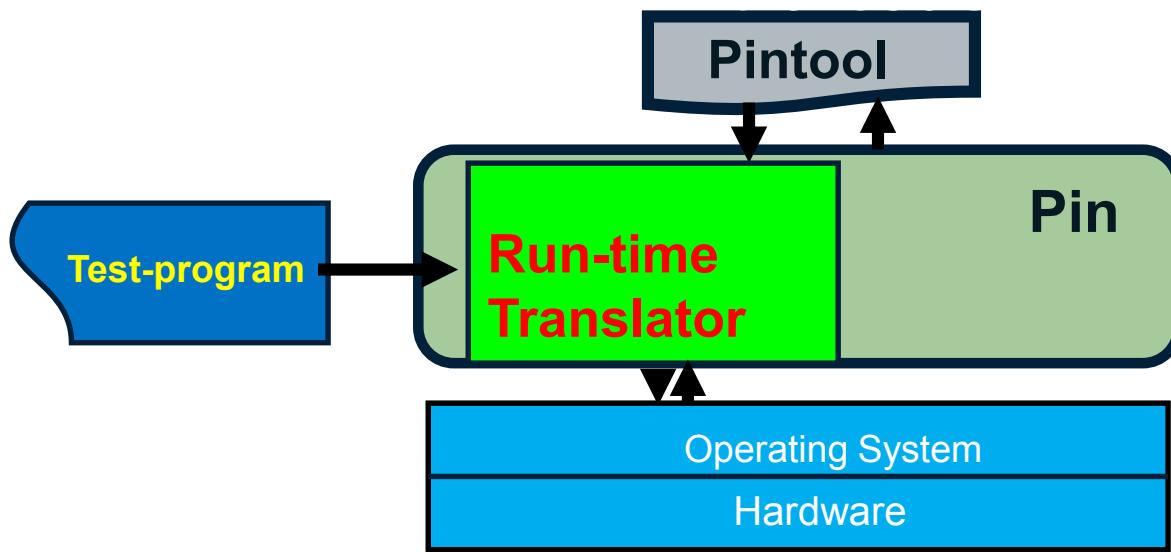
*To show that PinPlay is an **easy-to-use**,  
**flexible**, and **effective** framework for  
reproducible analysis of parallel  
programs.*

1. How PinPlay works
2. How to use it for reproducible analysis  
(e.g. debugging, dynamic program slicing, DCFG)

# Pin: A Tool for Writing Program Analysis Tools

```
sub    $0xff, %edx  
movl  0x8(%ebp), %eax  
jle   <L1>
```

```
counter++; print(IP)  
sub    $0xff, %edx  
counter++; print(EA)  
movl  0x8(%ebp), %eax  
counter++;print(br taken)  
jle   <L1>
```



```
$ pin -t pintool -- test-program
```

Normal output  
+ *Analysis output*

Pin: A Dynamic Instrumentation Framework from Intel  
<http://www.pintool.org>

# Schedule

9:00 -- 9:45 PinPlay and PinADX

9:45 – 10:15 DrDebug demo

10:15 – 10:30 Testimonial ( Milind)

10:30 – 11:00 Slicing demo

<11 – 11:20 Break>

11:20 – 11:50 DCFG (Chuck)

11:50 – 12:15 Maple (Cristiano)

12:15 – 12:30 Wrap-up + Q&A (all)

# The Need : Easier Analysis of Parallel Programs

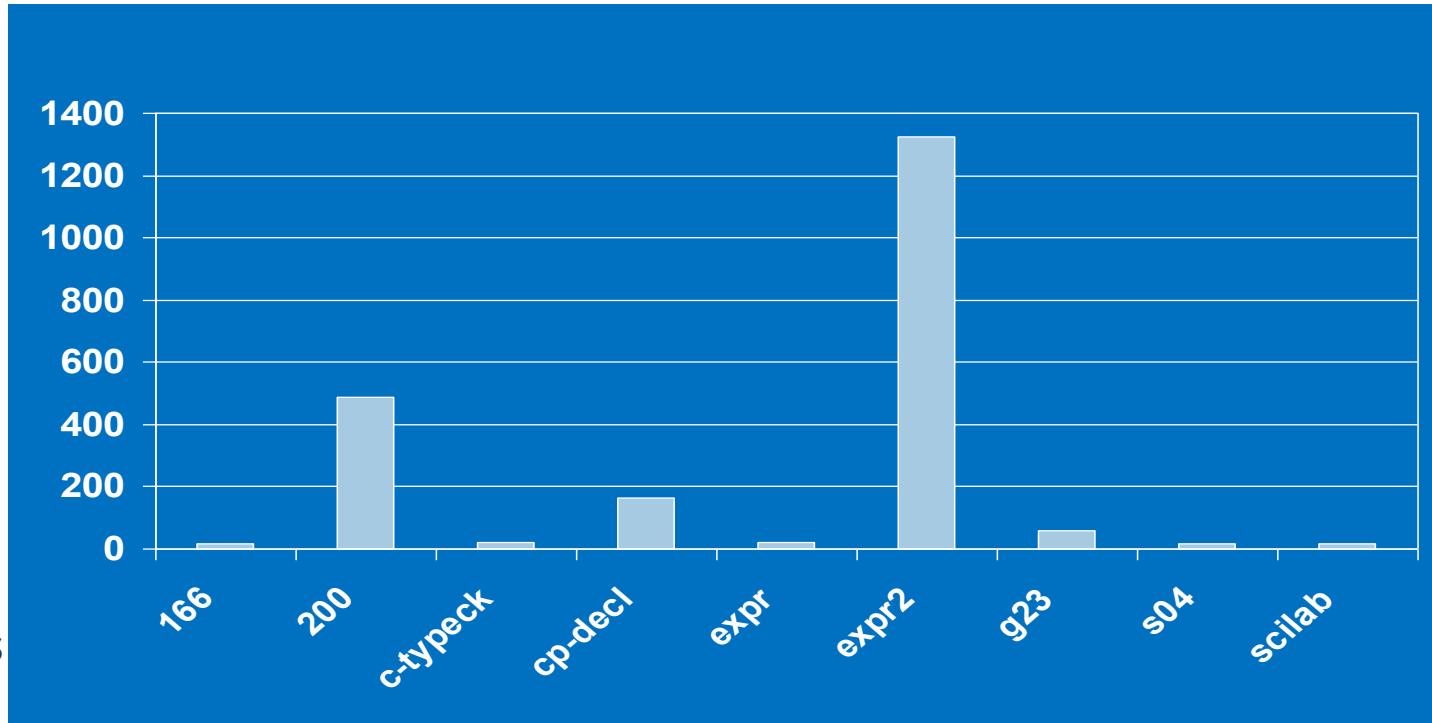
***Programmers need a way to deterministically analyze and debug parallel programs***

## Why?

Run-to-run variation → Chasing a moving target

# Run-to-run Control-flow Divergence in 403.gcc (SPEC2006)

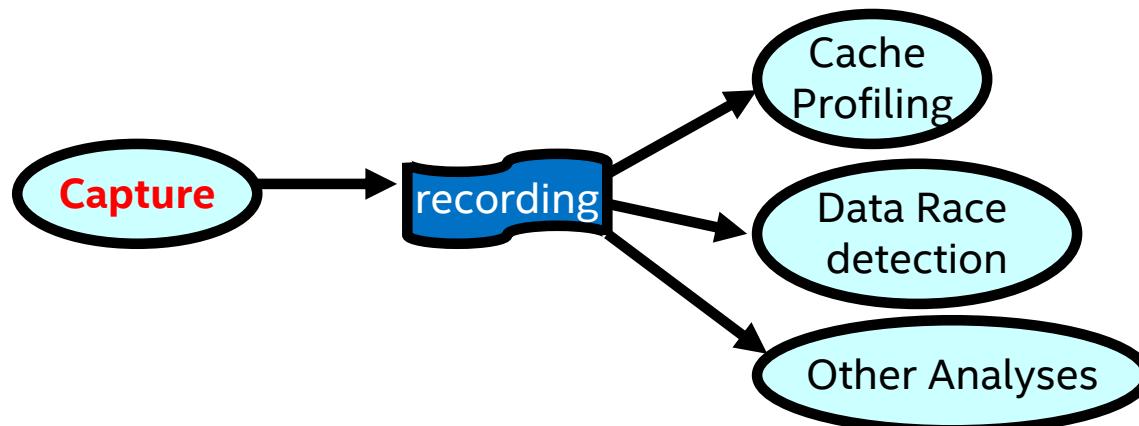
# of divergences between two runs  
8 'ref' inputs for gcc



Measured using an execution differencing tool based on PinPlay [Bert Maher]

**Even Single-threaded program runs are not repeatable!**

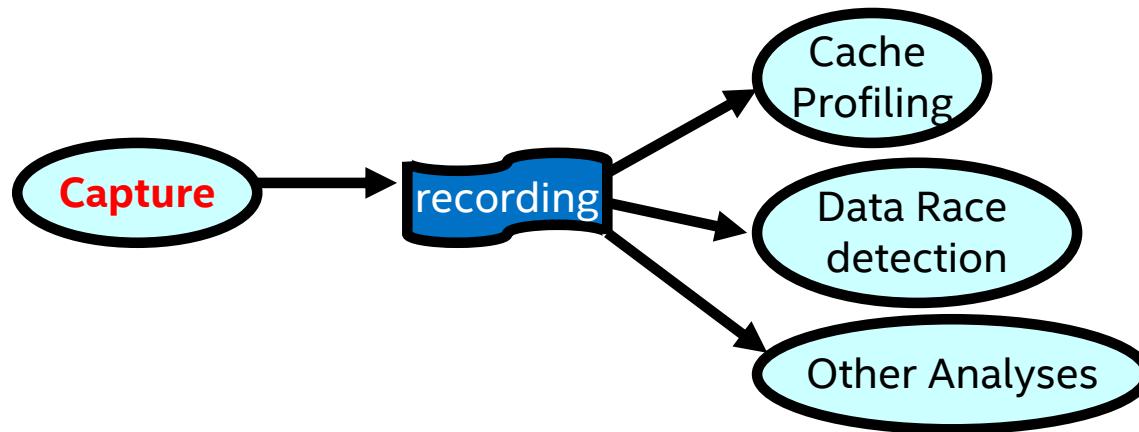
# The Solution : Record → Replay & Analyze



**Capture an execution and replay it deterministically with analysis**

# Grand Vision Capture once Analyze Multiple times! **Anywhere!**

Expensive analyses can be delayed till replay time with guaranteed repeatability



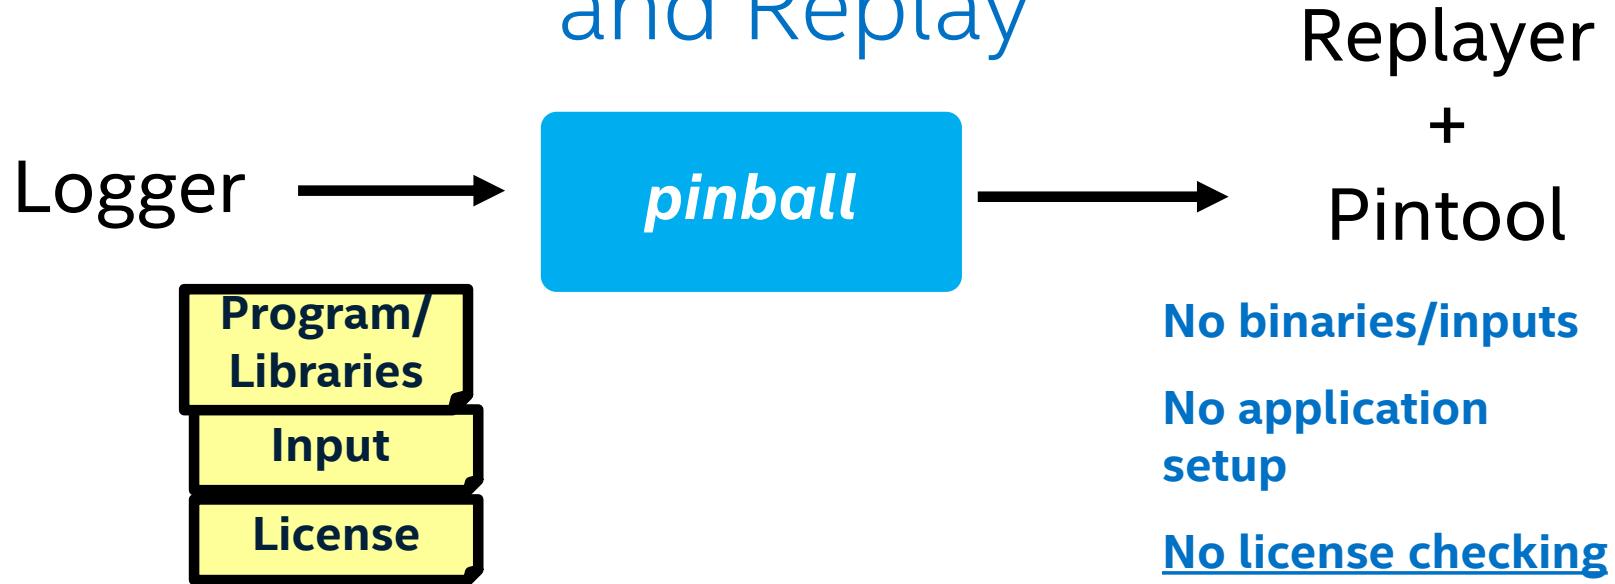
**Single Machine → NetBatch**

**2.4 Kernel → 2.6 Kernel**

**Windows → Linux**

**Customer site → Developer site**

# PinPlay : Software-based User-level Capture and Replay

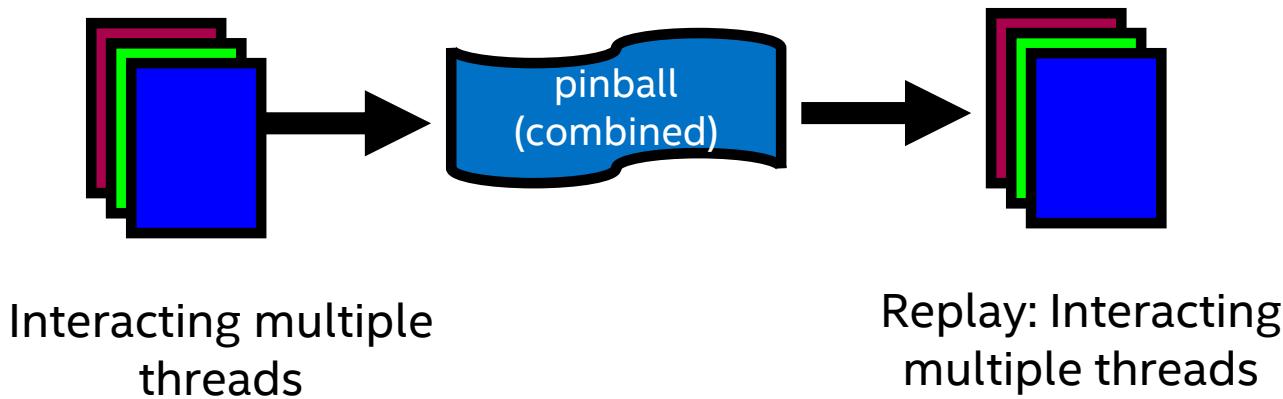


**Platforms :** Linux, Windows, Android, MacOS

**Upside :** It works! Large OpenMP / MPI programs, Oracle

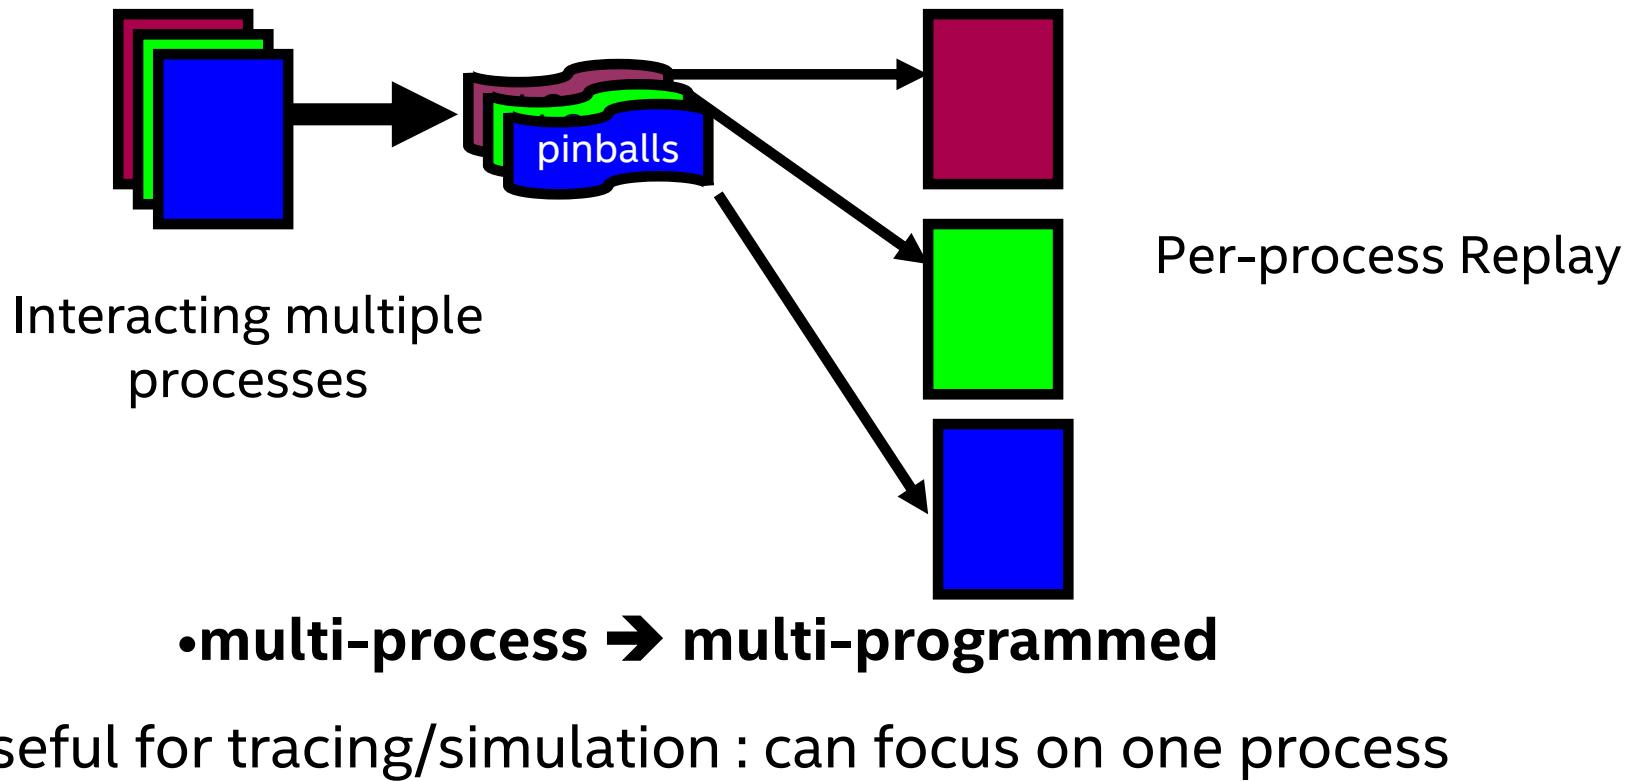
**Downside :** High run-time overhead: ~100-200X for capture → Cannot be turned on all the time

# Model 1: Parallel Capture : Parallel Replay For Multi-threaded Programs



**Useful for parallel analysis/debugging**

# Model 2: Parallel Capture : Isolated Replay For Multi-process Programs



# PinPlay : History & Acknowledgements

**Motivation** : Repeatable  
*PinPoints*

**Inspiration:** *BugNet* work from  
UC San Diego

**Break-through** : Automatic  
system call side-effect analysis

**Initial implementation** for  
deterministic multi-threaded  
simulation

**Further development +  
Support** : Windows port,  
Android port, multi-threaded  
region recording, tracing...

**Why?** “*PinPoints out of order*”  
27/55 SPEC2006 benchmarks!

Satish Narayanasamy, Giles Pokam,  
Prof. Brad Calder [ISCA 2005]

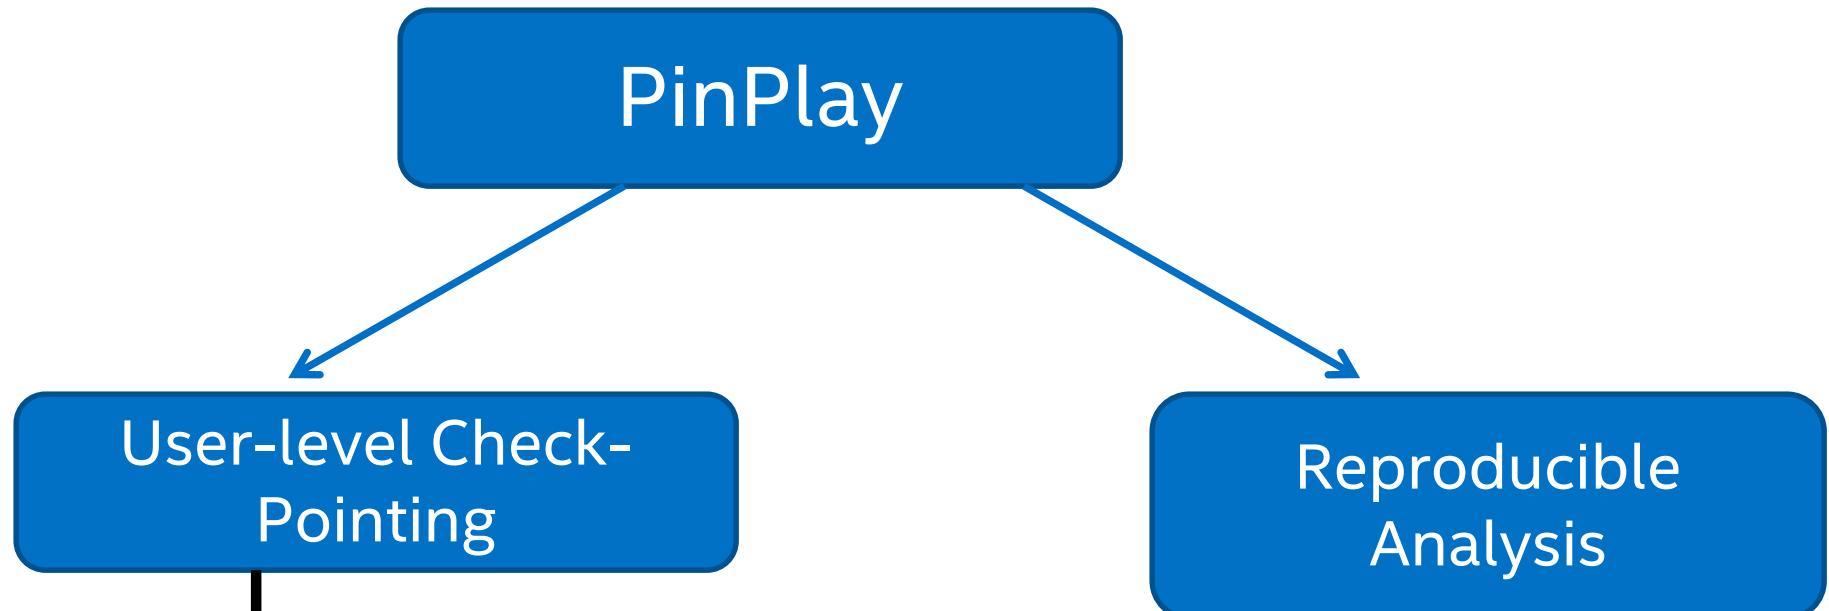
Satish Narayanasamy, Cristiano  
Pereira... [SIGMETRICS 2006]

Cristiano Pereira (Ph. D. thesis)

Jim Cownie, Ady Tal, Ariel Slonim,  
Michael Gorin, Michael Berezalsky,  
Tevi Devor, Omer Mor, Roman  
Zendel, Mack Stallcup, Cristiano  
Pereira, Harish Patil, **Pin team**

**Sponsors:** Geoff Lowney, Robert  
Cohn, Moshe Bach, Sion Berkowits,  
Nafta Shalev

# PinPlay Applications at a glance



1. Simulation region selection (*PinPoints*)
2. Dynamic program slicing (**UC Riverside**)
3. Replay-based debugging (*DrDebug*)
4. Dynamic control-flow graph generation
5. < Your analysis here >

# PinPlay basics

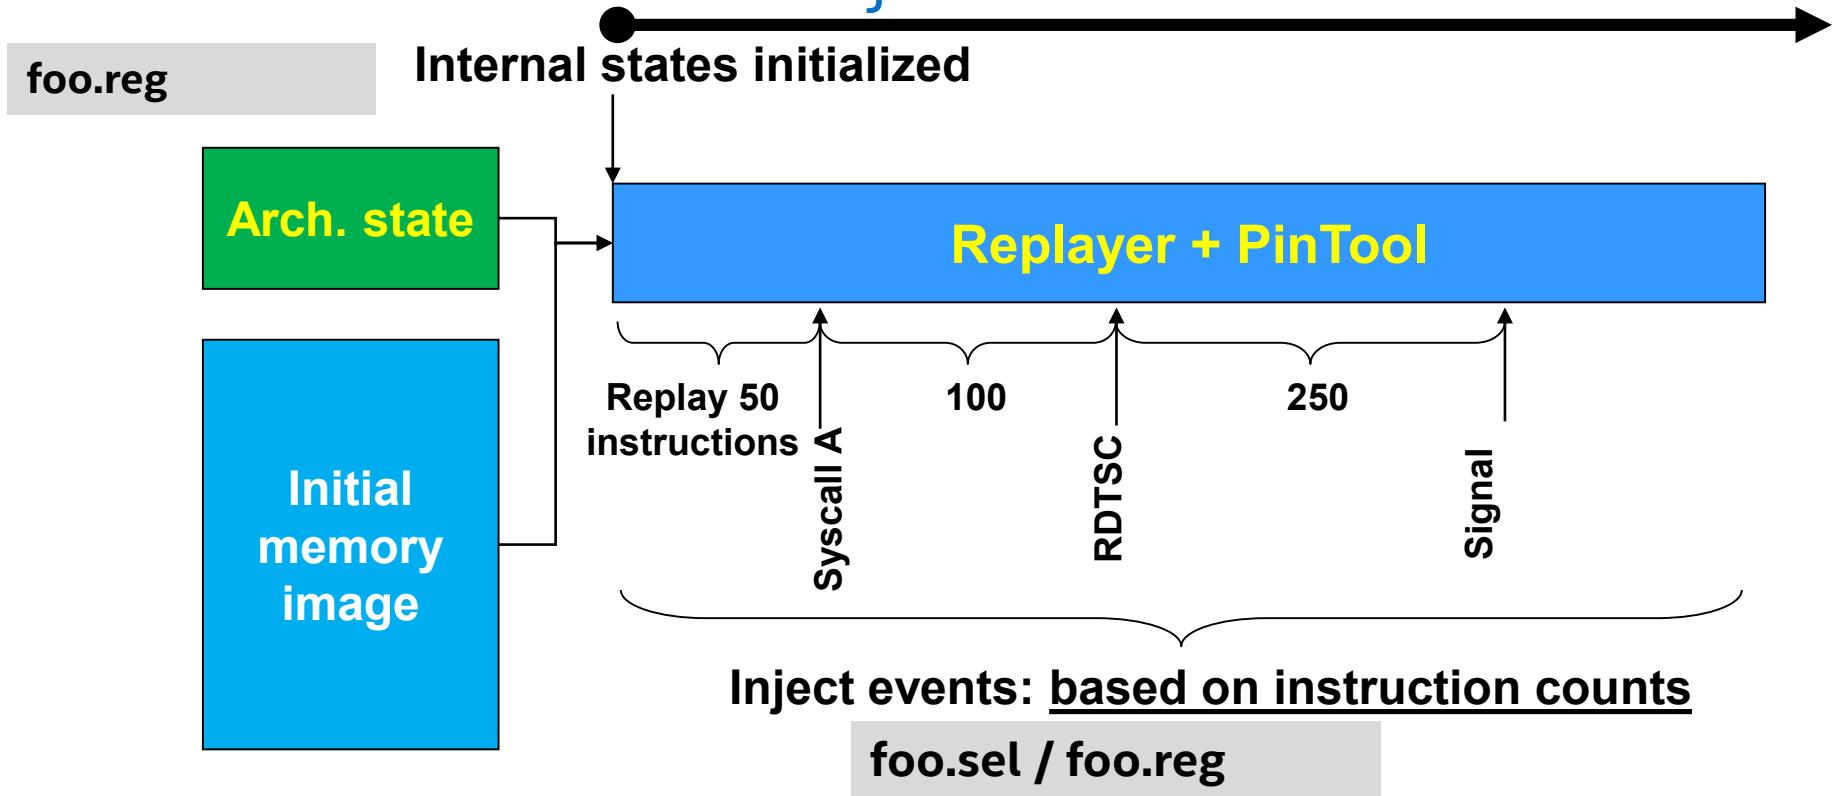
# What makes single-threaded program runs non-repeatable?

1. Initial **stack** location changes (by kernel); changes to the location of **dynamically allocated memory** (heap)
2. **Shared library load location** changes (randomized address space);
3. Program binary/shared-library **code** changes (change over time or from machine to machine)
4. Processor-specific instruction behavior changes:  
**CPUID, RDTSC**
5. **Signals:** arbitrary memory/register/control-flow changes
6. **Un-initialized memory** location reads
7. **System call** behavior changes (depending on workdirectory, environment variables. e.g: *gettimeofday()* *uname()*)

} Memory logging

**PinPlay efficiently handles all these**

# Pinball (ST) = Initial memory/register + injections



**foo.text**

- **System calls** : skipped by injecting next rip/ memory changed
- **CPUID, RDTSC** : affected registers injected
- **Signals/Callbacks** : New register state injected

# Multi-threaded programs: Additional source of non-repeatability

- Access **order** to shared-memory locations may change from run-to-run
  - Un-protected shared memory accesses
  - Accesses to locks/mutexes

**Why:** Threads may have different rates of progress across runs

- OS scheduling, machine load, machine configuration

# Multi-threaded programs: Repeating Shared Memory Access Order

**Logging:** Record (a subset of) RAW, WAW, WAR dependences

**Replaying:** Satisfy dependences; make threads wait if needed

***How to detect RAW/WAW/WAR dependences?***

***Simulate a directory-based cached coherence protocol in software***

**Netzer optimization : avoids logging implied dependencies**  
**SPECOMP : Only 0.22% of all dependencies actually logged!**

# Pinball (MT) : Pinball (ST) + Thread-dependencies

**foo.reg (per-thread)**



**foo.text**

**Application Memory (common)**

**foo.reg (per-thread)**

**foo.sel (per-thread)**

**Event injection works only if same behavior  
(same instruction counts) is guaranteed  
during replay**

**foo.race (per-thread)**

[T1] 2 T2 2  
[T1] 3 T2 3

[T2] 5 T4 1

Thread T2 cannot execute instruction 5  
until T4 executes instruction 1

Thread T1 cannot execute instruction 2  
until T2 executes instruction 2

# Sources of Slow-down

## 1. All programs (single-threaded or parallel) :

**Logger:** Instrument loads/stores (memory logging)

**Replayer:** Instrument loads : Restore memory at the right “time”

## 2. Parallel programs:

**Logger:** All memory instruction analyses guarded by locks to prevent changes to memory during analysis

If parallel replay desired:

**Logger:** Emulate directory-based cache-coherence protocol in software

**Replayer:** Obey logged shared-memory dependences; make threads wait

**Logging more expensive than replay**

**Parallel replay more expensive than isolated replay**

# Special Case : Single-threaded Programs

## Fast Logging and Replay possible

### **Fast Logging (not implemented yet):**

- Ad-hoc system-call side-effect analysis and logging  
(only practical for Linux)

### **Fast Replay (implemented):**

- Restore memory locations “lump-sum” at system calls
  - Replay at Pin-only speed!  
(for single-threaded programs without signals: e.g.  
SPEC2006)

# Summary: Slowdown & Disk Usage

| Application                                   | Instruction Count (avg) | Pinball size (average) | Logger Slowdown : X native (avg) | Replayer Slowdown: X-native (avg)                          |
|---|-------------------------|------------------------|----------------------------------|--|
| SPEC2006 (55 runs) single-threaded            | 924 billion             | 39MB                   | 98 X                             | 11 X<br><u><a href="#">[1.4 X with “-replay:fast”]</a></u> |
| PARSEC (8 runs) : multi – threaded:<br>4T—25T | 844 billion             | 3.9GB                  | 197 X                            | 37 X   |
| SPECOMP 2001 : 4T                             | 307 billion             | 91 MB                  | 117X                             | 25X  |

- Measured with PinPlay 2.0 ‘pin -t pinplay-driver.so’
- Some PARSEC benchmarks read large files which increases pinball size substantially

# The PinPlay kit



# Download from <http://www.pinplay.org>

```
pinplay-drdebug-2.2-pldi2015-pin-2.14-71313-gcc.4.4.7-linux
|-- doc
|   '-- html
|-- extras
|   '-- pinplay
|       |-- PinPoints
|       |   '-- bin
|       '-- bin
|           |-- ia32
|           '-- intel64
|       '-- examples
|           '-- tests
|       '-- include
|       '-- include-ext
|       '-- lib
|           |-- ia32
|           '-- intel64
|       '-- lib-ext
|           |-- ia32
|           '-- intel64
|       '-- scripts
|           '-- pinplay.H
|               '-- libpinplay.a
|               '-- libslicing.a
|                   |-- gdb_record
|                   |-- gdb_replay
|                   '-- record
|                   '-- relog
|                   '-- replay
|           '-- dcfg
|               |-- bin
|               '-- doc
|               '-- examples
|               '-- include
|               '-- lib
```

The diagram illustrates the file structure of the `pinplay-drdebug-2.2-pldi2015-pin-2.14-71313-gcc.4.4.7-linux` distribution. Red arrows highlight the following correspondences between the main tree and the expanded view:

- `bin` points to `StaticAnalysis`, `pinplay-branch-predictor.so`, and `pinplay-driver.so`.
- `examples` points to `pinplay-branch-predictor.cpp`, `pinplay-debugger-shell.cpp`, and `pinplay-driver.cpp`.
- `include` points to `pinplay.H`.
- `lib` points to `libpinplay.a` and `libslicing.a`.
- `scripts` points to `gdb_record`, `gdb_replay`, `record`, `relog`, and `replay`.
- `dcfg` points to `bin`, `doc`, `examples`, `include`, and `lib`.

# Enabling a Pintool for PinPlay

```
#include "pinplay.H"
```

```
PINPLAY_ENGINE pinplay_engine;
```

```
KNOB<BOOL>KnobReplayer(KNOB_MODE_WRITEONCE, KNOB_FAMILY,  
                           KNOB_REPLY_NAME, "0", "Replay a pinball");  
KNOB<BOOL>KnobLogger(KNOB_MODE_WRITEONCE, KNOB_FAMILY,  
                           KNOB_LOG_NAME, "0", "Create a pinball");
```

```
pinplay_engine.Activate(argc, argv, KnobLogger, KnobReplayer);
```

**Link** in *libpinplay.a*, *libzlib.a*, *libbz2.a*,  
**\$ (CONTROLLERLIB)**

## Restriction:

- PinTool shouldn't change application control flow

# Example: pinplay-branch-predictor.cpp

```
#define KNOB_LOG_NAME "log"
#define KNOB_REPLAY_NAME "replay"
#define KNOB_FAMILY "pintool;pinplay-driver"

PINPLAY_ENGINE pinplay_engine;

KNOB_COMMENT pinplay_driver_knob_family(KNOB_FAMILY, "PinPlay Driver Knobs");

KNOB<BOOL>KnobReplayer(KNOB_MODE_WRITEONCE, KNOB_FAMILY,
                         KNOB_REPLAY_NAME, "0", "Replay a pinball");
KNOB<BOOL>KnobLogger(KNOB_MODE_WRITEONCE, KNOB_FAMILY,
                      KNOB_LOG_NAME, "0", "Create a pinball");

int main(int argc, char *argv[])
{
    if( PIN_Init(argc,argv) )
    {
        return Usage();
    }

    outfile = new ofstream(KnobStatFileName.Value().c_str());
    bimodal.Activate(KnobPhases, outfile);

    pinplay_engine.Activate(argc, argv, KnobLogger, KnobReplayer);

    PIN_AddThreadStartFunction(threadCreated, reinterpret_cast<void *>(0));

    PIN_StartProgram();
}
```

# PinPlay-enabled PinTools : 3 Modes

## 1. Regular Analysis mode

```
$ pin -t pintool -- test-program
```

Normal output  
+ Analysis  
output

## 2. Logging Mode

```
$ pin -t  
pintool -log -log:basename pinball/foo -- test-program
```

## 3. Replay Mode

```
$ pin -t pintool
```

```
-replay -replay:basename pinball/foo
```

**pinball**

**nullapp**

# Example: pinplay-branch-predictor.so

```
% pin -t $PIN_ROOT/extras/pinplay/bin/intel64/pinplay-  
branch-predictor.so -- hello
```

Creates “**bimodal.out**”

```
%pin -t pinplay-branch-predictor.so -log -log:basename  
pinball/foo hello
```

Creates “**bimodal.out**” and “**pinball/foo\***”

```
%pin -xyzzy -reserve_memory pinball/foo.address -replay  
-replay:basename pinball/foo --  
$PIN_KIT/extras/pinplay/intel64/bin/nullapp
```

Creates “**bimodal.out**”

# Using \$PIN\_ROOT/extras/pinplay/scripts

## :Recording (uses pinplay-driver.so):

```
pinplay-VirtualBox:~/tests/hello> which record  
/home/pinplay/PinPlay/latest/extras/pinplay/scripts//record
```

```
% record --help  
Usage: record.py [options] -- binary args  
        or  
        record.py [options] --pid PID
```

Create a recording (pinball). There are two modes:

- 1) Give command line of a binary to record
- 2) Give the PID of a running process

```
pinplay-VirtualBox:~/tests/hello> record --pintool $PIN_ROOT/ex  
tras/pinplay/bin/intel64/pinplay-branch-predictor.so --pinball  
pinball/foo -- hello
```

\* Developed by *Mack Stallcup*

# Using \$PIN\_ROOT/extras/pinplay/scripts: Replaying (uses pinplay-driver.so)

```
pinplay-VirtualBox:~/tests/hello> which replay  
/home/pinplay/PinPlay/latest/extras/pinplay/scripts//replay
```

```
% replay --help  
Usage: replay.py [options] -- pinball  
  
Replay a recording (pinball).
```

```
pinplay-VirtualBox:~/tests/hello> replay --pintool $PIN_ROOT/ex  
tras/pinplay/bin/intel64/pinplay-branch-predictor.so -- pinball  
/foo_0  
0 added by 'record'
```

# Adding your own file to a pinball

## pinplay.H

```
//register to start/stop logging callbacks  
VOID RegisterRegionStart(PINPLAY_HANDLER handler, VOID* args);  
VOID RegisterRegionStop(PINPLAY_HANDLER handler, VOID* args);
```

## PinPlay-enabled Pin-tool

### Register callbacks

```
// Region start/stop callbacks from PinPlay when logging.  
pinplay_engine.RegisterRegionStart(RegionStart, 0);  
pinplay_engine.RegisterRegionStop(RegionStop, 0);
```

### START handler

```
// Start a region when logging.  
VOID RegionStart(VOID* v)  
{  
    THREADID pin_tid = PIN_ThreadId();  
    string region_basename = pinplay_engine->GetRegionBaseName(pin_tid);  
    // Open file you are adding -- use 'region_basename' as prefix.  
    // $region_basename.myext
```

### STOP handler

```
// Stop a region when logging.  
VOID RegionStop(VOID* v)  
{  
    // close your file $region_basename.myext
```

# Replay Debugging Foundation: PinADX



Developed by **Greg Lueck** (Intel Corporation)

Slides from **Tevi Devor** (Intel Corporation)

# Transparent debugging, and extending the debugger

## Transparently debug the application while it is running on Pin + Pin Tool

- **PinADX: Customizable Debugging with Dynamic Instrumentation** (Presented at CGO 2012)

## Use Pin Tool to enhance/extend the debugger capabilities

- Watchpoint: Is order of magnitude faster when implemented using Pin Tool
- Which branch is branching to address 0
  - Easy to write a Pin Tool that implements this

# Debug Application while Running Pin

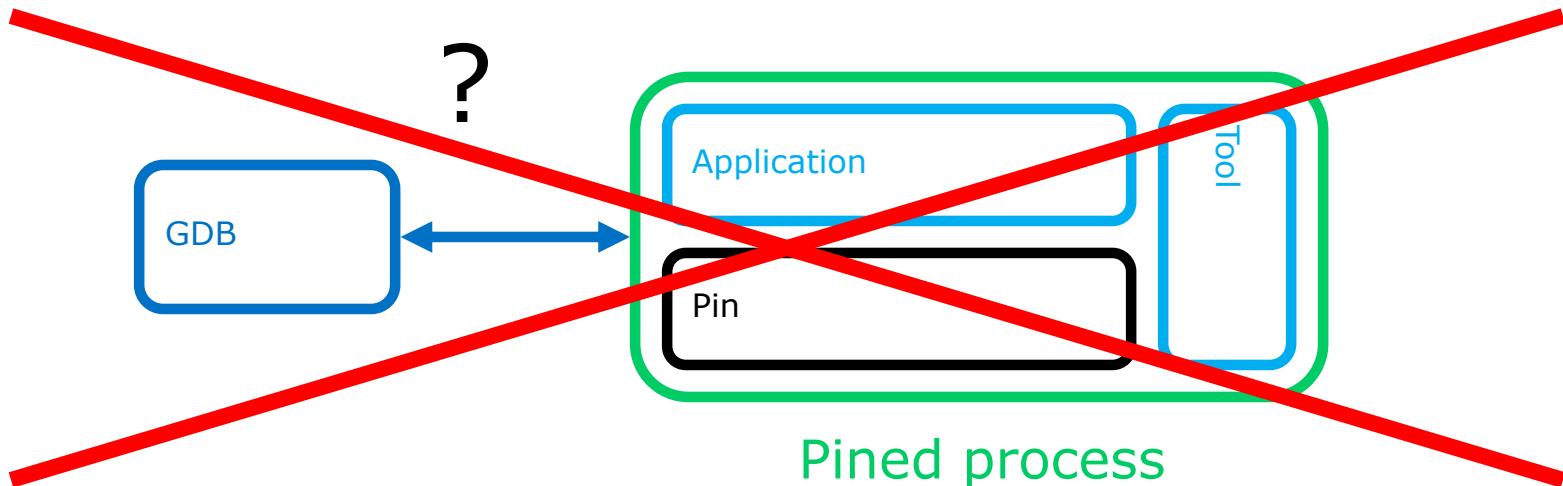
Useful for Pin-based emulators

- User can debug application while emulating

Provide advanced debugging features with Pin:

- Stack monitoring features
- Buffer overrun detection
- Reverse debugging
- Write your own debugger extension via Pin

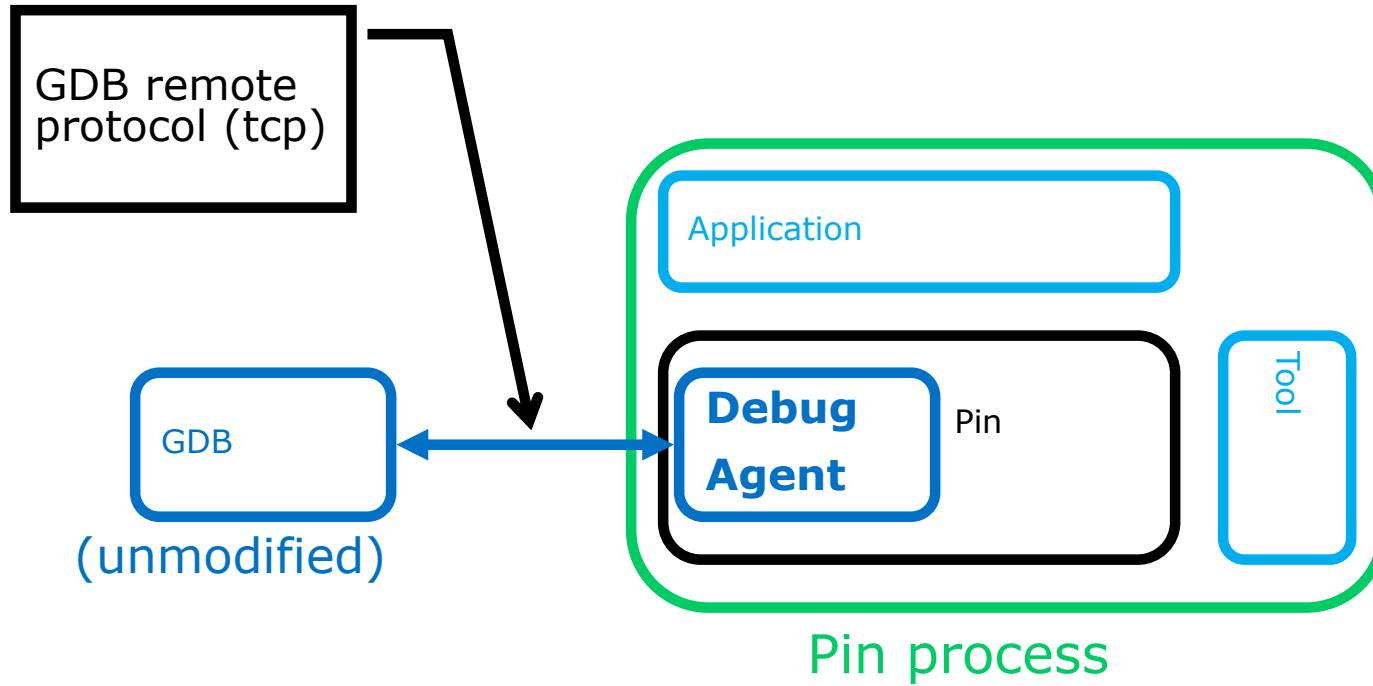
# Naïve Solution Won't Work



Why can't we just debug normally?

- Debugger sees Pin state, not application state
- Pin recompiles application code
- Instructions wrong, registers wrong, PC wrong, ...

# Pin Debugger Interface



GDB debugs application (not Pin itself)

Leverage GDB remote protocol ABI

# Debug the Application with Pin

## 1. Run Pin with -appdebug

```
$ pin -appdebug -t tool.so -- ./application
Application stopped until continued from debugger.
Start GDB, then issue this command at the (gdb) prompt:
target remote :1234
```

## 2. Start GDB, enter “target remote ...”

```
$ gdb ./application
(gdb) target remote :1234
```

## 3. Set breakpoints, etc. Continue with “cont”

```
(gdb) break main
(gdb) cont
```

*[DrDebug scripts : gdb\_record and gdb\_replay :  
Single window solution]*

# Extending the Debugger

Normal debugging with Pin useful but limited

Extending the debugger:

- Add GDB commands via a Pin tool
- Stop at “semantic breakpoint” via instrumentation

Use the “monitor” keyword for implementing custom commands

```
// Debugger interpreter, to process debugger commands.  
//  
PIN_AddDebugInterpreter(DebugInterpreter, this);
```

[DrDebug : See *extras/pinplay/examples/pinplay-debugger-shell.cpp*]

[DrDebug: uses “pin.py” to provide “pin <extended-command>”]

# Stack Debugger Pintool

```
$ pin -appdebug -t stack-debugger.so -- ./app
```

```
$ gdb ./app
```

```
(gdb) target remote :1234
```

```
(gdb) monitor stackbreak 4000
```

Break when thread uses 4000 stack bytes.

```
(gdb) cont
```

Thread uses 4004 stack bytes.

```
[...]
```

```
(gdb) monitor stats
```

Maximum stack usage: 8560 bytes.

Commands implemented  
in Pintool

# Stack-Debugger Instrumentation

Thread Start:

**StackBase = %esp;**

**MaxStack = 0;**

[...]

sub \$0x60, %esp  
**size = StackBase - %esp;**

**if (size > MaxStack) MaxStack = size;**

**if (size > StackLimit) TriggerBreakpoint();**

cmp %esi, %edx

jle <L1>

After each stack-changing instruction



# ManualExamples/stack-debugger.cpp

```
VOID Instruction(INS ins, VOID *)           instrumentation routine
{
    if (INS_RegWontain(ins, REG_STACK_PTR) )
    {
        IPOINT where = (INS_HasFallThrough(ins)) ?
            IPOINT_AFTER : IPOINT_TAKEN_BRANCH;
        INS_InsertCall(ins, where, (AFUNPTR)OnStackChange,
            IARG_REG_VALUE, REG_STACK_PTR,
            IARG_THREAD_ID, IARG_CONST_CONTEXT, IARG_END);
    }
}

VOID OnStackChange(ADDRINT sp, THREADID tid, CONTEXT *ctxt)
{
    size_t size = StackBase - sp;
    if (size > StackMax) StackMax = size;           analysis routine
    if (size > StackLimit) {
        ostringstream os;
        os << "Thread uses " << size << " stack bytes.";
        PIN_ApplicationBreakpoint(ctxt, tid, FALSE, os.str());
    }
}
```

Insert only after instructions  
that write to %esp

# ManualExamples/stack-debugger.cpp

```
VOID Instruction(INS ins, VOID *)           instrumentation routine
{
    if (INS_RegWContain(ins, REG_STACK_PTR))
    {
        IPOINT where = (INS_HasFallThrough(ins)) ?
            IPOINTER_After : IPOINTER_TAKEN_BRANCH;
        INS_InsertCall(ins, where, (AFUNPTR)OnStackChange,
            IARG_REG_VALUE, REG_STACK_PTR,
            IARG_THREAD_ID, IARG_CONST_CONTEXT, IARG_END);
    }
}

VOID OnStackChange(ADDRINT sp, THREADID tid, CONTEXT *ctxt)
{
    size_t size = StackBase - sp;
    if (size > StackMax) StackMax = size;           analysis routine
    if (size > StackLimit) {
        ostringstream os;
        os << "Thread uses " << size << " stack bytes.";
        PIN_ApplicationBreakpoint(ctxt, tid, FALSE, os.str());
    }
}
```

Triggers debugger  
breakpoint

# stack-debugger.cpp

Hooks the GDB “monitor” command. E.g.:  
(gdb) monitor stats  
(gdb) monitor stackbreak 4000

```
int main() {  
    [...]  
    PIN_AddDebugInterpreter(HandleDebugCommand, 0);  
}  
  
BOOL HandleDebugCommand(const string &cmd, string *result) {  
    if (cmd == "stats")  
    {  
        ostringstream os;  
        os << "Maximum stack usage: " << StackMax << " bytes.\n";  
        *result = os.str();  
        return TRUE;  
    }  
    else if (cmd.find("stackbreak ") == 0)  
    {  
        StackLimit = /* parse limit */;  
        ostringstream os;  
        os << "Break when thread uses " << StackLimit << " stack bytes.";  
        *result = os.str();  
        return TRUE;  
    }  
    return FALSE; // Unknown command  
}
```

[DrDebug : See extras/pinplay/examples/pinplay-debugger-shell.cpp]

# ManualExamples/stack-debugger.cpp

```
int main() {  
    [...]  
    PIN_AddDebugInterpreter(HandleDebugCommand, 0);  
}  
  
BOOL HandleDebugCommand(const string &cmd, string *result) {  
    if (cmd == "stats")  
    {  
        ostringstream os;  
        os << "Maximum stack usage: "  
        *result = os.str();  
        return TRUE;  
    }  
    else if (cmd.find("stackbreak ") == 0)  
    {  
        StackLimit = /* parse limit */;  
        ostringstream os;  
        os << "Break when thread uses " << StackLimit << " stack bytes.";  
        *result = os.str();  
        return TRUE;  
    }  
    return FALSE; // Unknown command  
}
```

Receives text after "monitor"

This string written to GDB session

[DrDebug : See extras/pinplay/examples/pinplay-debugger-shell.cpp]

# Other Debugger Tools

Breakpoint on buffer overrun

Debug from a recorded log file

Reverse debugging from a recording

Design your own custom debugger tool

## DrDebug : Debugging, Analysis, and Replay Toolkit

See `extras/pinplay/examples/pinplay-debugger-shell.cpp`

# Replay Debugging: DrDebug kit

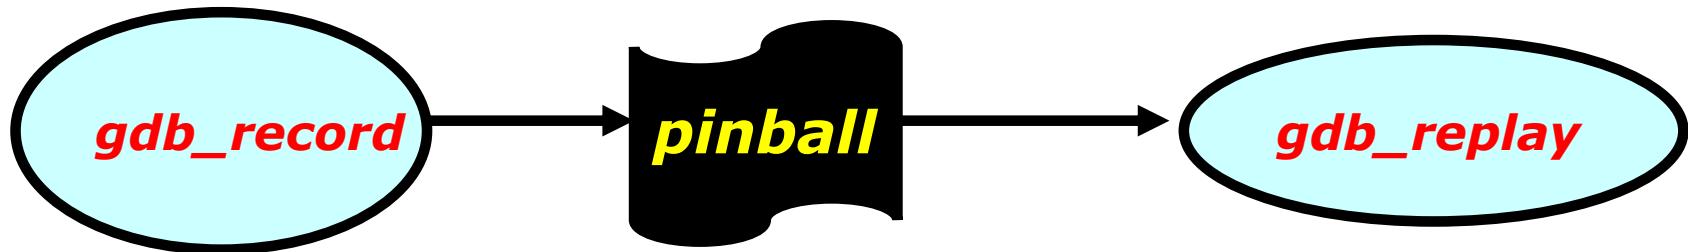




1. What is DrDebug?
2. Why DrDebug?
3. How DrDebug works?
4. Demo

# What is DrDebug

- Pin-based GDB extensions  
*gdb\_record* and *gdb\_replay*
- Supports replay-based debugging



- Brings the power of Pin's dynamic analysis to GDB  
pintools → Extended GDB commands  
(fast conditional breakpoints, `printf` debugging..)

**DrDebug/PinPlay kit @ [www.drdebug.org](http://www.drdebug.org)**  
**Eclipse GUI: CDT plugin PinPlay→Dynamic-Slicing**

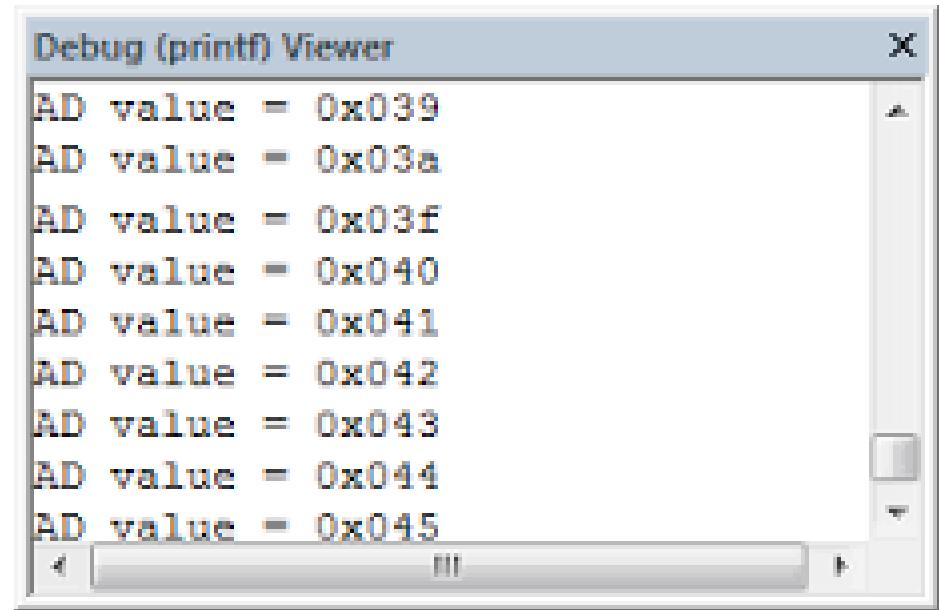
# Why DrDebug?

*“Because programmers are people too!”*

# How people debug?

## 1. *printf* debugging

**What:** Print values/messages at points of interest



### Challenges:

1. Requires source/re-compilation
2. Values/control-flow may change across runs

# Tracing Variables with DrDebug

- Use Pin to instrument points of interest and print values
- No source changes/re-compilation necessary
- Debugging with replay → Values do not change across sessions

```
% gdb_replay -- failing.pinball/log_0 bread-demo
```

```
(gdb) pin trace order at 189
monitor trace [ %rbp + -8 ] 8 at 0x40171d #order:<189>
Tracepoint #1: trace memory [%rbp offset -8 ] length 8
at 0x40171d #order:<189>
```

```
Breakpoint 1, 0x00002aaac0884160 in _exit ()
(gdb) pin trace print to order.txt
```

```
% head order.txt
0x000000000040171d: [$rbp + -8] = 0x551bb0 #order:<189>
0x000000000040171d: [$rbp + -8] = 0x550550 #order:<189>
0x000000000040171d: [$rbp + -8] = 0x5492b0 #order:<189>
0x000000000040171d: [$rbp + -8] = 0x5486f0 #order:<189>
0x000000000040171d: [$rbp + -8] = 0x547350 #order:<189>
```

# How people debug?

## 2. Cyclic debugging



0. Make a hypothesis about the bug  
while (not (bug-root-caused))  
do

1. Fast forward till the buggy region
2. Proceed slowly, examine state
3. Hit the bug; more state examination

done



# Cyclic Debugging Challenges

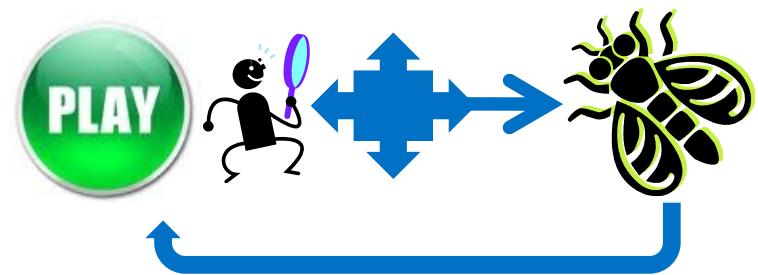
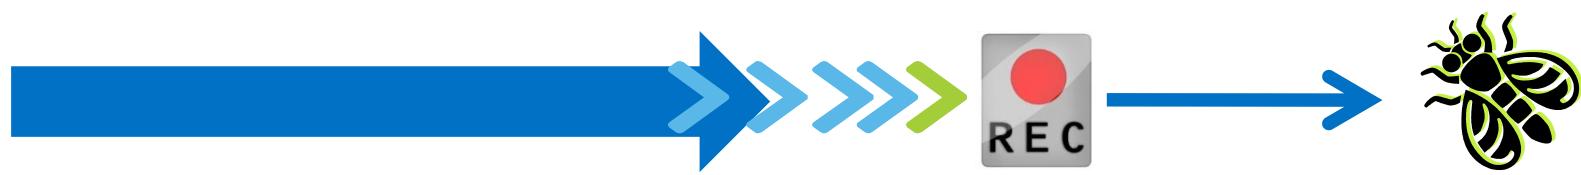


1. **Takes too long to reach bugs at times** : every debug cycle is long
2. **Many things change on each debug cycle iteration** : stack, heap, signals, interaction with outside world (interactive programs), network IO, thread schedule....
3. **Some bugs are hard to repeat in general and also under debugger**

1. Takes too long to reach bugs at times : Each debug cycle iteration is long

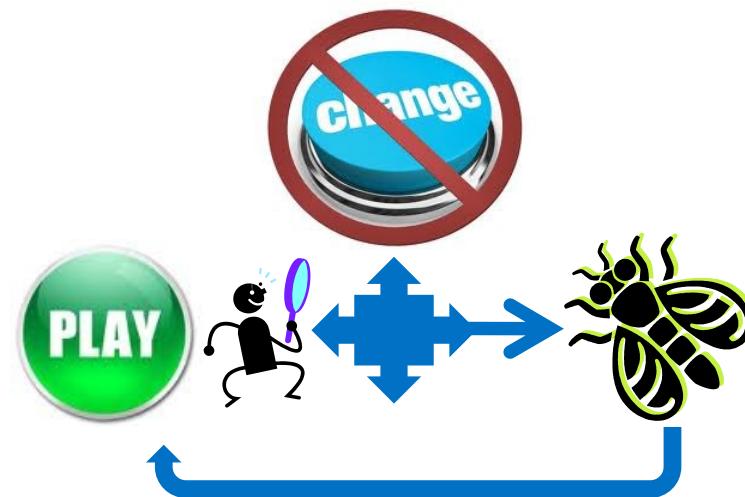
**Once:** Go to buggy region, hit record → capture bug

**Iteratively** : Start at the buggy region not the beginning of the program



2. Many things change on each debug cycle iteration  
: stack, heap, thread schedule, system call  
results...

DrDebug guarantees stack/heap/other non-deterministic events  
are repeatable



### 3. Some bugs are hard to repeat in general and also under debugger



***Once captured, the bug never escapes***

*How do I get the bug to occur under the debugger?*

Expose the bug by changing thread schedules

What we did: Interfaced with Maple bug exposing tool

**(See Maple presentation later)**

# So, what is the cost? Record/Replay Everything → High Overhead

**Why:** No HW/OS dependencies; Guaranteed repeatability

| Program/input          | Logging Slowdown<br>(X native) | Replayer Slowdown<br>(X native) |
|------------------------|--------------------------------|---------------------------------|
| SPEC2006/ref (ST)      | 98X                            | 11X                             |
| PARSEC/'native'(4-25T) | 197X                           | 37X                             |



Not needed for debugging!

# Debug Determinism

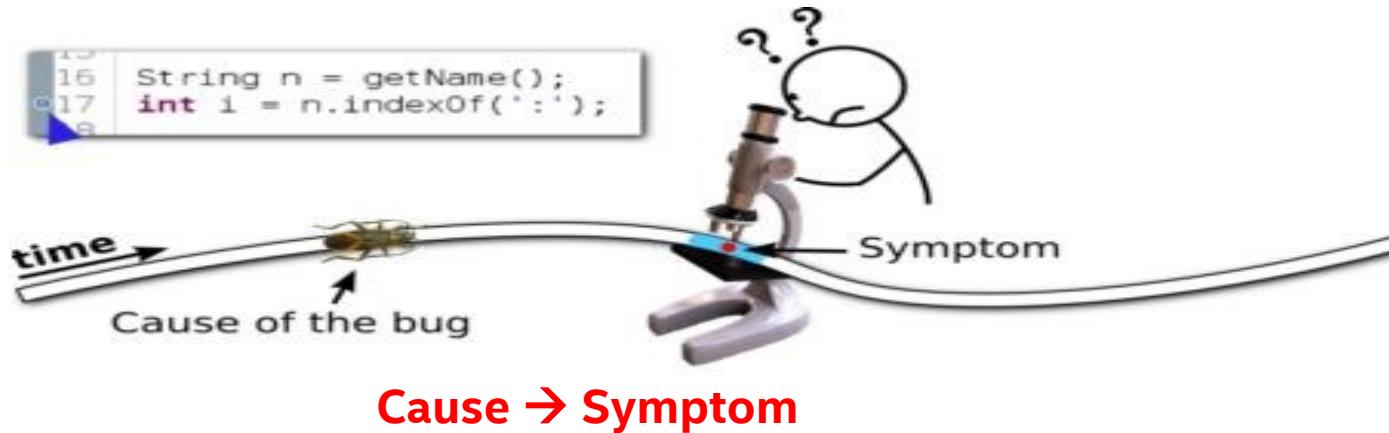
To be useful for debugging, should at a minimum reproduce the original failure and root cause.



**Skip:** To interesting region  
(before cause)

**Cause → Symptom**

# So, what's the cost, again? Depends on “Buggy Region” length

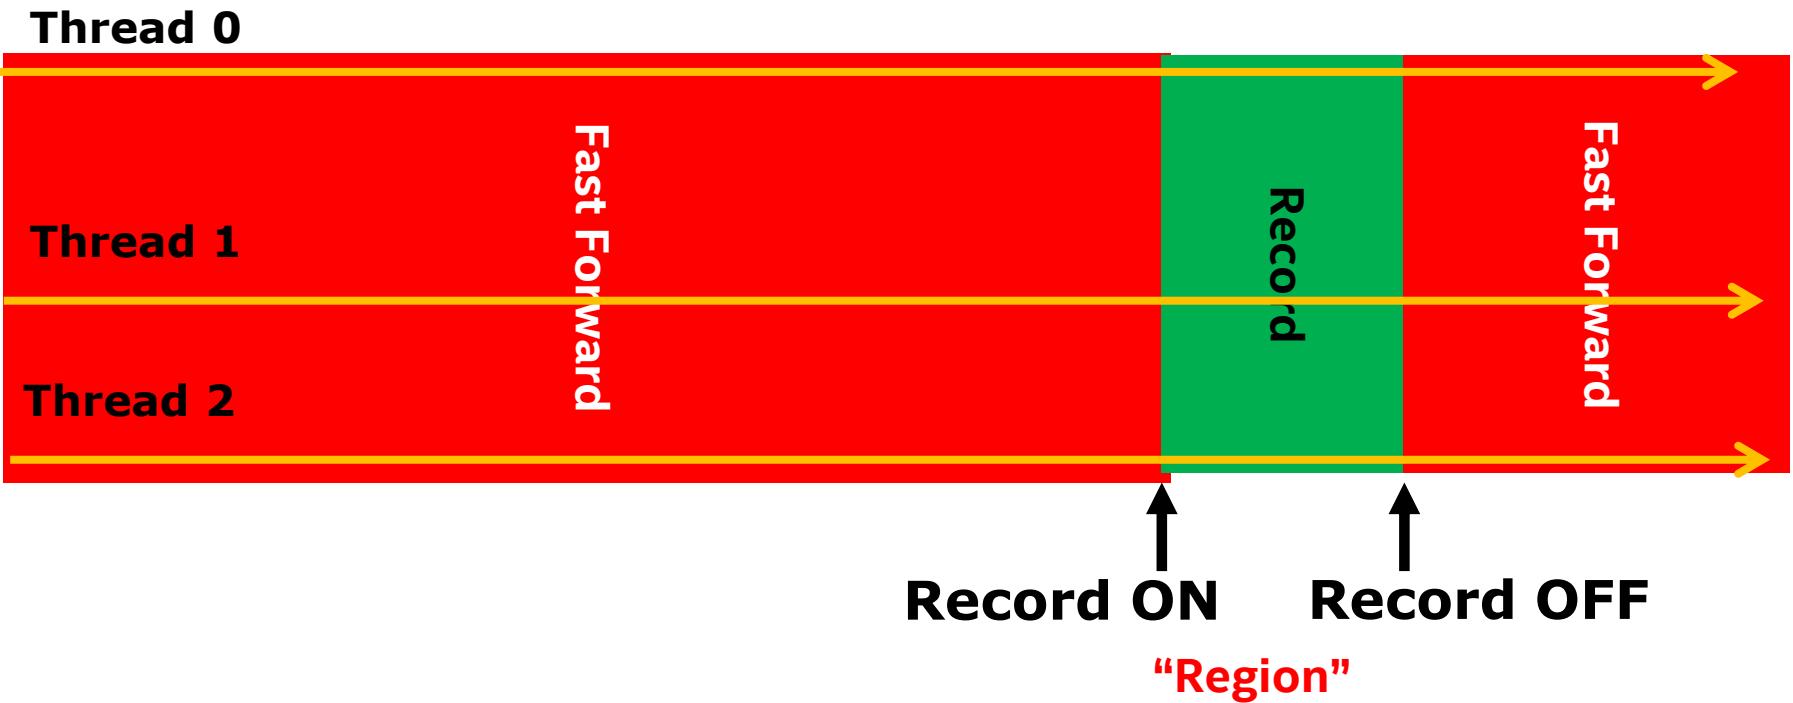


## Typical “buggy region” lengths:

1. ~10 million instructions (*study of 13 buggy open source programs*)
2. 8 million(Mozilla), 76K(Aget) : (our testing CGO2014)

# Terminology: What is a “region”?

Points in execution where recording “starts” and “ends”.



**Buggy Region:** captures **Root Cause → Symptom**

# Example region record/replay times

PARSEC (native input)– 4 Threaded runs

**Average times in seconds for 8 programs**

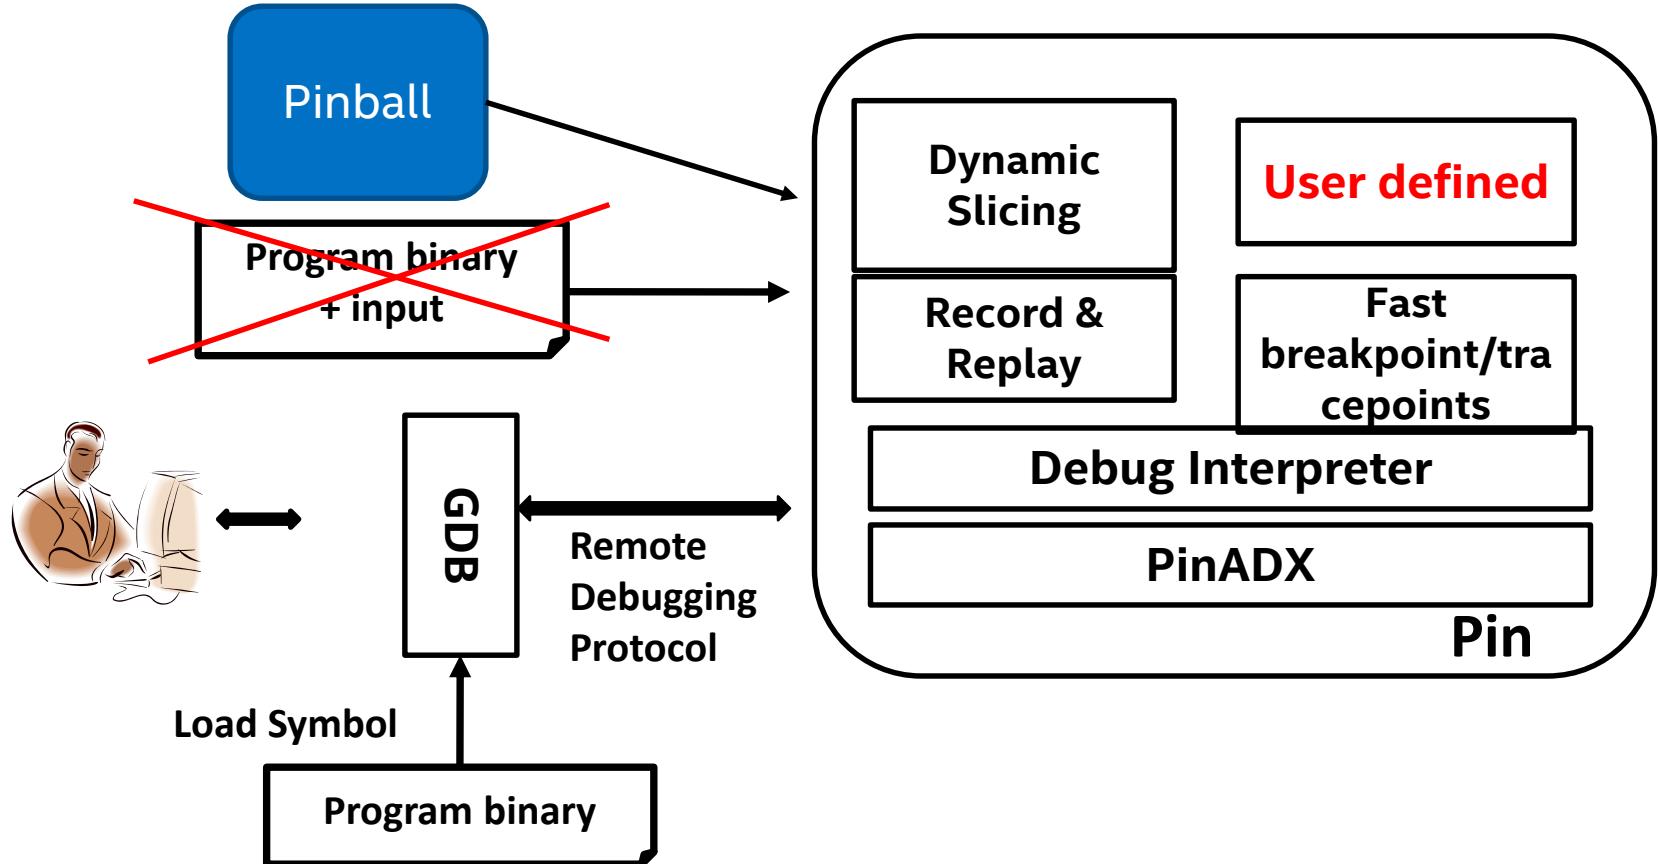
**Region length in ‘instruction count’**

**( region triggered in one thread : all threads captured)**

## Cause → Symptom

| Region length   | Region record time | Region replay time |
|-----------------|--------------------|--------------------|
| 4 x 10 million  | 10 seconds         | 10 seconds         |
| 4 x 100 million | 29 seconds         | 27 seconds         |
| 4 x 500 million | 75 seconds         | 43 seconds         |
| 4 x 1 billion   | 125 seconds        | 55 seconds         |

# DrDebug: Linux Setup



<http://www.drdebug.org>

# Test Program : bread-demo

Processes **10000 orders** with a pool of **4 threads**

Prints “**WRONG!**” if check fails

## Master

Initialize Bakery Database

Spawn 4 worker threads

Gather sub-totals: compute TOTAL

Compute MYTOTAL

**TOTAL ==  
MYTOTAL?**

NO

“**WRONG!**”

## Worker X 4

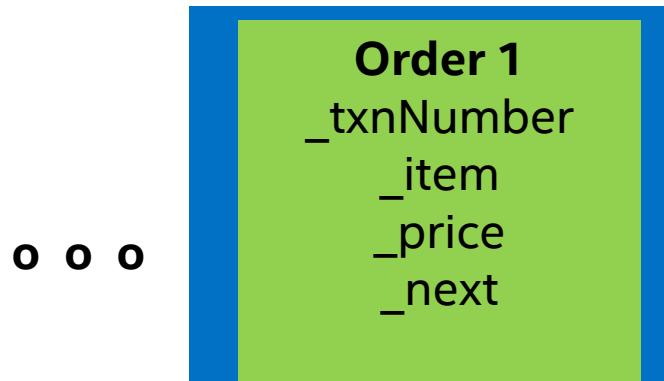
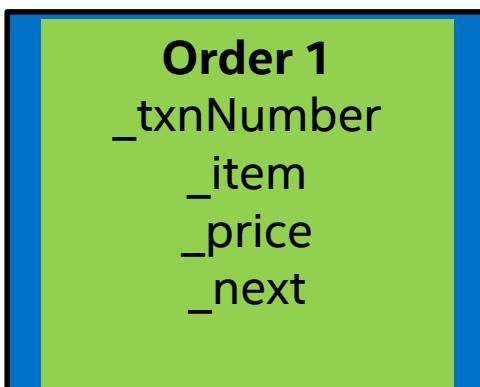
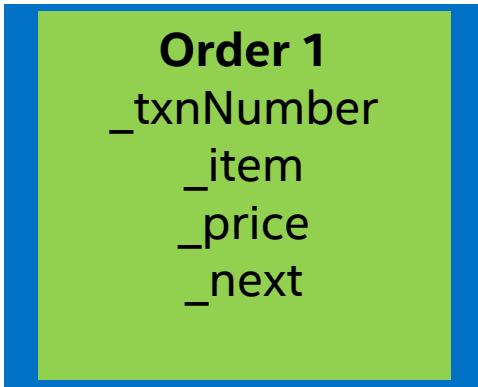
GetLock(bakery)

Grab an “order”

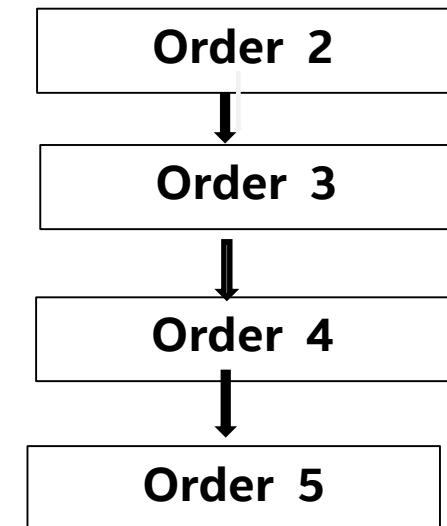
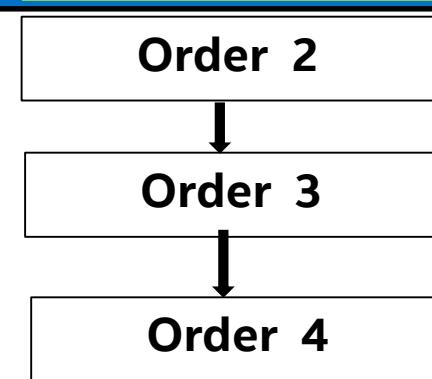
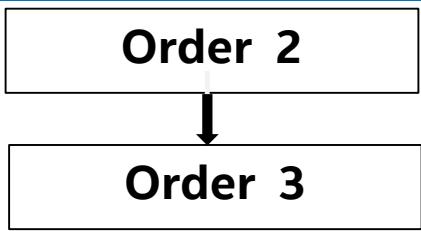
Unlock

Update SUBTOTAL

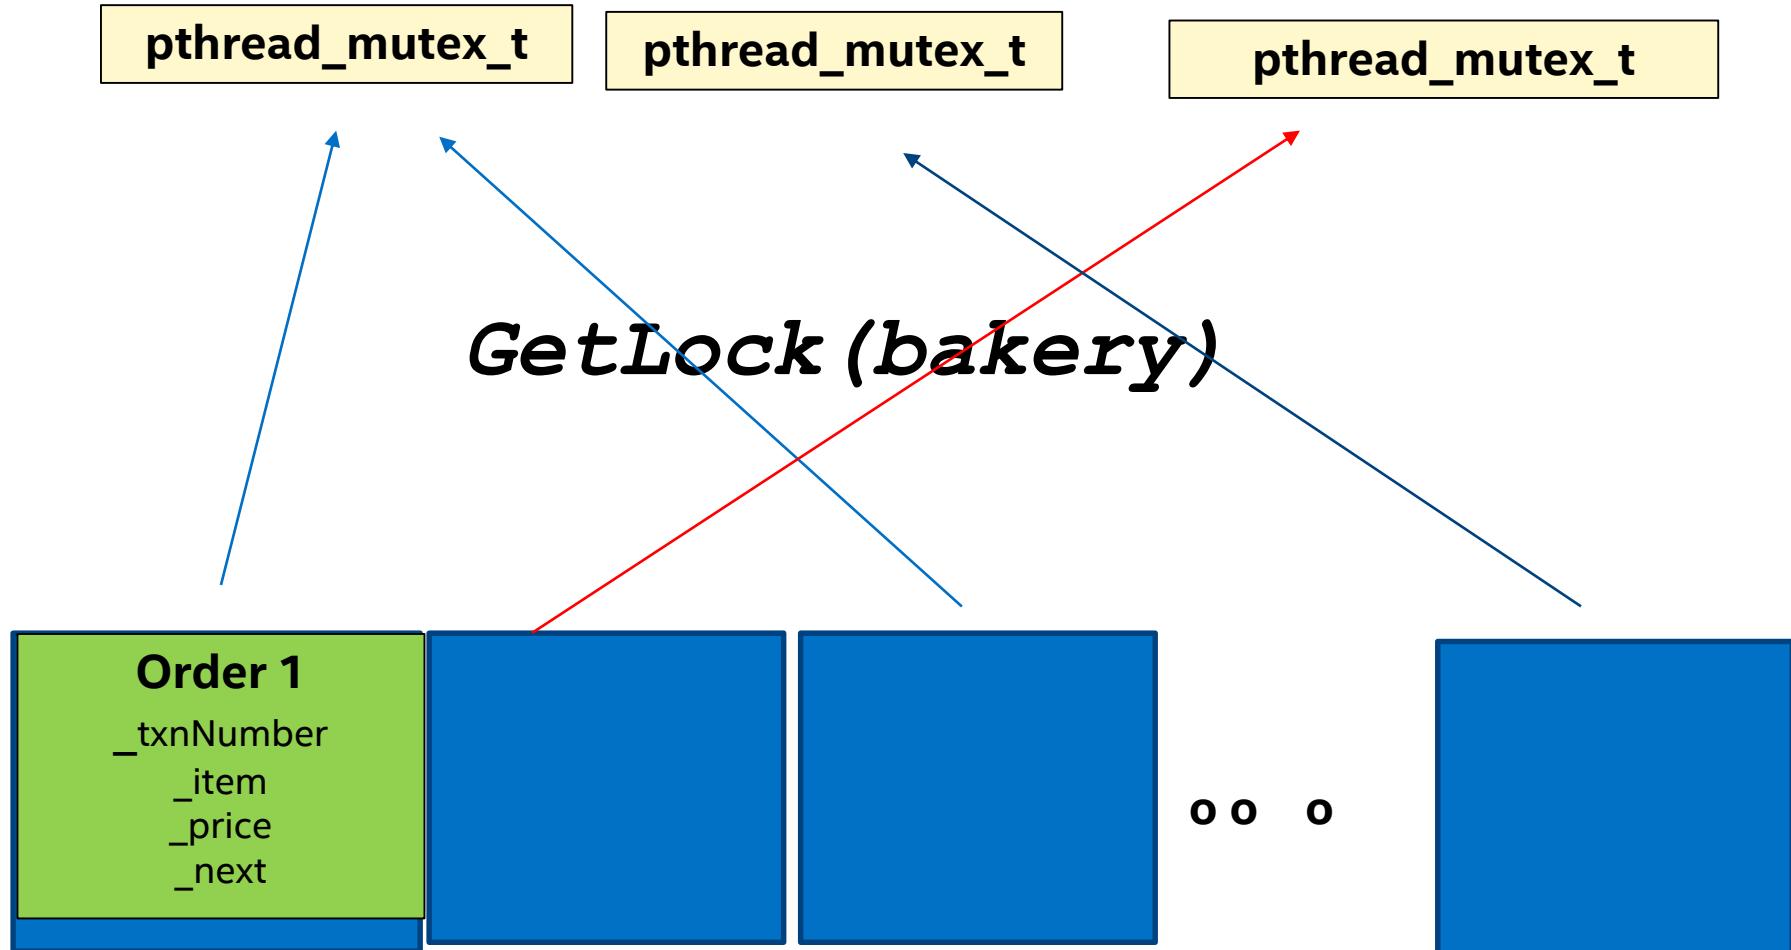
# ORDER \*Bakeries[NUMBER\_BAKERIES (= 500)]



o o o



# Pthread\_mutex\_t Locks[NUMBER\_LOCKS (= 8)]



**Need a unique lock per bakery**

```
static void *WorkerMain(void *v)
{
    // Loop over each bakery's queue, processing one transaction from
    // each queue before moving to the next queue.
    //
    for (;;)
    {
        bakery = (bakery + 1) % NUMBER_BAKERIES;

        // Get the lock for this bakery's queue.
        //

        pthread_mutex_t *lock = GetLock(bakery);
        pthread_mutex_lock(lock);

        // Get one transaction from this queue.
        //
        ORDER *order = Bakeries[bakery];

Line 189 Bakeries[bakery] = order->_next;


        pthread_mutex_unlock(lock);

        // Process the transaction.
        //
        worker->breadOrders[order->_item]++;
        worker->totalRevenue += order->_price;
    }
    // We're done when there are no transactions left on any bakery's queue.
    return 0;
}
```

# Demo: Instructions

```
pinplay-VirtualBox:~/DrDebug/BreadDemo.pldi> cat HOWTO.txt
Record a bug:
-----
gdb_record bread-demo
b 113
c
pin record on
c
q
# if the session above shows a "WRONG" message, you have captured a buggy
# execution; else repeat the session.
Replay a buggy execution:
-----
replay pinball/log_0

gdb+replay a buggy execution:
-----
gdb_replay pinball/log_0 bread-demo
pin trace order at 189
b _exit
c
pin trace print to order.txt
q

Find order being processed twice/multiple times:
-----
% sort.sh order.txt
capture order value XXXX that appears multiple times
```

# Demo: Instructions (continued)

```
gdb+replay a buggy execution:  
-----  
gdb_replay pinball/log_0 bread-demo  
pin break 189 if order == XXXX  
c  
p order  
info threads  
# see which thread is processing order XXXX  
c  
# second processing of the same order  
p order  
info threads  
# see which thread is processing order XXXX  
# at this point we know which two threads are processing the same order.  
# if both the threads are still at line 189, you can switch between threads  
thread m  
print lock  
thread j  
print lock  
#two threads have different locks, why?  
Continue exploring....
```

# Recording a “*buggy region*” with gdb\_record

```
Terminal  
pinplay-VirtualBox:~/DrDebug/BreadDemo.pldi> gdb_record bread-demo  
GNU gdb (Ubuntu 7.7.1-0ubuntu5~14.04.2) 7.7.1  
Copyright (C) 2014 Free Software Foundation, Inc.  
  
-----  
(gdb) b 113  
Breakpoint 1 at 0x4011fd: file bread-demo.cpp, line 113.  
(gdb) c  
-----
```

```
Breakpoint 1, main (argc=1, argv=0x7fff812181a8) at bread-demo.cpp:113  
113          Go = true;  
(gdb) pin record on  
monitor record on  
Started recording region number 0
```

```
(gdb) c  
Continuing.
```

```
Total revenue: $21948.42 (WRONG!)  
[Inferior 1 (Remote target) exited with code 01]  
(gdb)
```

# Replay + debugging with gdb\_replay

```
pinplay-VirtualBox:~/DrDebug/BreadDemo.pldi> replay pinball/log_0

Bread sales for WW04
-----
white      : 1720
wheat      : 1644
bagel      : 1728
baguette   : 1674
croissant  : 1606
challah    : 1630

Total revenue: $21948.42 (WRONG!)
```

```
pinplay-VirtualBox:~/DrDebug/BreadDemo.pldi> gdb_replay pinball/log_0 bread-demo

GNU gdb (Ubuntu 7.7.1-0ubuntu5~14.04.2) 7.7.1
```

```
118          double revenue = 0.0;
(gdb) pin trace order at 189
```

```
(gdb) b _exit
Breakpoint 1 at 0x7f74587242d0: _exit. (2 locations)
(gdb) c
```

# Replay + debugging with gdb\_replay (contd.)

```
Total revenue: $21948.42 (WRONG!)

Breakpoint 1, __GI__exit (status=status@entry=1)
  at ../sysdeps/unix/sysv/linux/_exit.c:28
28      ../sysdeps/unix/sysv/linux/_exit.c: No such file or directory.
(gdb) pin trace print to order.txt
monitor trace print to order.txt
(gdb) c
Continuing.
[Inferior 1 (Remote target) exited normally]
(gdb) █
```

```
pinplay-VirtualBox:~/DrDebug/BreadDemo.pldi> wc -l order.txt
10002 order.txt
pinplay-VirtualBox:~/DrDebug/BreadDemo.pldi> sort.sh order.txt
 2 0x0000000000401533: [$rbp + -8] = 0x1f45550 #order:<189>
 2 0x0000000000401533: [$rbp + -8] = 0x1f28e90 #order:<189>
 1 0x0000000000401533: [$rbp + -8] = 0x1f4c1f0 #order:<189>
 1 0x0000000000401533: [$rbp + -8] = 0x1f4c1d0 #order:<189>
 1 0x0000000000401533: [$rbp + -8] = 0x1f4c1b0 #order:<189>
```

# Replay + debugging with gdb\_replay (contd.)

```
pinplay-VirtualBox:~/DrDebug/BreadDemo.pldi> gdb_replay pinball/log_0 bread-demo  
GNU gdb (Ubuntu 7.7.1-0ubuntu5~14.04.2) 7.7.1
```

```
(gdb) pin break 189 if order == 0x1f45550  
monitor break at 0x401533 if [ %rbp + -8 ] 8 == 0x1f45550 #order:<189>  
Breakpoint #1: break at 0x401533 if [$rbp offset -8 ] length 8 == 0x1f45550 #  
order:<189>  
(gdb) c  
Continuing.
```

```
(gdb) p order  
$1 = (ORDER *) 0x1f45550  
(gdb) p bakery  
$2 = 4  
(gdb) info threads  
[New Thread 9011]  
[New Thread 9010]  
[New Thread 9009]  
 Id  Target Id          Frame  
 5   Thread 9009      0x000000000040153e in WorkerMain (v=0x1f4c210)  
     at bread-demo.cpp:189  
 4   Thread 9010      WorkerMain (v=0x1f4c238) at bread-demo.cpp:171  
 3   Thread 9011      WorkerMain (v=0x1f4c288) at bread-demo.cpp:171  
* 2   Thread 9012      WorkerMain (v=0x1f4c260) at bread-demo.cpp:189  
 1   Thread 9001      0x00007f7458f7966b in pthread_join (  
     threadid=140137651791616, thread_return=0x0) at pthread_join.c:92  
(gdb) p lock  
$3 = (pthread_mutex_t *) 0x6043f0 <Locks+240>
```

# Replay + debugging with gdb\_replay (contd.)

```
(gdb) c
Continuing.
Triggered breakpoint #1: break  at 0x401533 if  [$rbp offset -8 ] length 8 == 0x
1f45550 #order:<189>
Program received signal SIGTRAP, Trace/breakpoint trap.
[Switching to Thread 9010]
WorkerMain (v=0x1f4c238) at bread-demo.cpp:189
189          Bakeries[bakery] = order->_next;
(gdb) p order
$4 = (ORDER *) 0x1f45550
(gdb) p bakery
$5 = 4
(gdb) info threads
  Id  Target Id        Frame
  5  Thread 9009      0x000000000040153e in WorkerMain (v=0x1f4c210)
    at bread-demo.cpp:189
* 4  Thread 9010      WorkerMain (v=0x1f4c238) at bread-demo.cpp:189
  3  Thread 9011      0x00000000000400f00 in pthread_mutex_lock@plt ()
  2  Thread 9012      WorkerMain (v=0x1f4c260) at bread-demo.cpp:189
  1  Thread 9001      0x00007f7458f7966b in pthread_join (
    threadid=140137651791616, thread_return=0x0) at pthread_join.c:92
```

# Demo : Takeaway

- We could do multiple `gdb_replay` sessions and gather more information each time
- This was feasible because
  1. the same buggy schedule was reproduced in each session and
  2. pointer values ('order' values) remained the same across sessions
- That is the power of replay debugging –  
**bug once captured never escapes!**

# Replay Debugging: User Testimonial

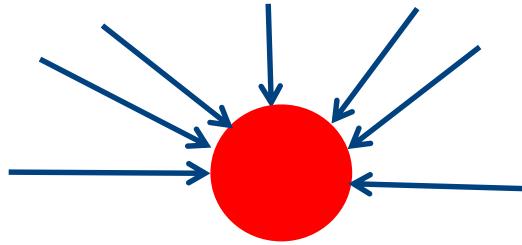


**Milind Chabbi : Rice University**

# A Quick Review of Locks

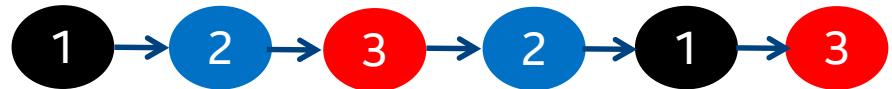
- [BAD] Centralized locks (Test&Set, Test&Test&Set, Backoff)

- ✗ Remote traffic
  - ✗ Poor scaling
  - ✗ Unfair



- [GOOD] Queuing locks (MCS, CLH)

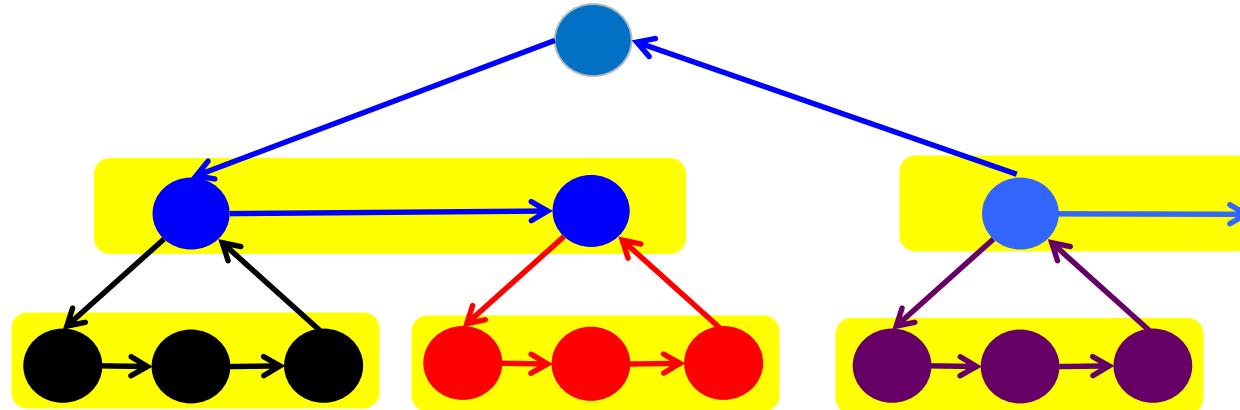
- ✓ Local spinning
  - ✓ Superior scaling
  - ✓ FIFO fairness



**✗ Indiscriminate data movement on NUMA architectures**

# Hierarchical MCS (HMCS): A NUMA-Aware Lock

- Always pass the lock to a thread in the nearest NUMA domain
  - Reuse the temporal locality of reference
- Bound the local passing to ensure starvation freedom
- Result: ~70x higher throughput than MCS lock on 4096-core SGI UV 1000

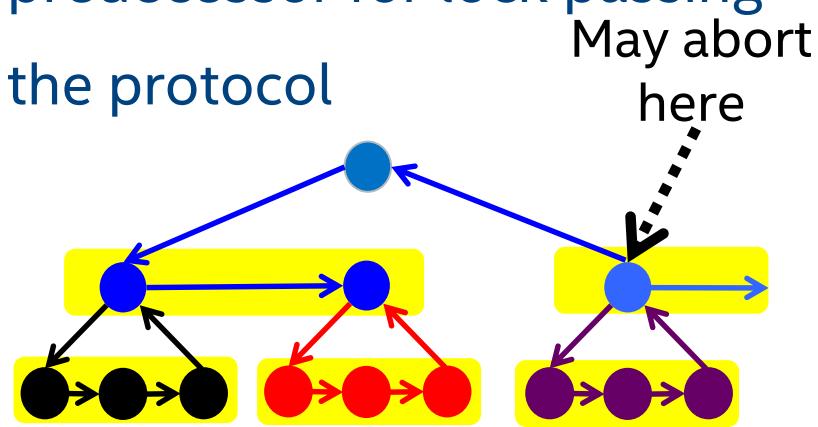


[PPoPP'15] High performance locks for multi-level NUMA systems. Chabbi et al.

# A Sophisticated Abortable HMCS Lock

Allow a waiting thread to **timeout** and leave the protocol

- Complications due to
  - State-full nature of the protocol and massive asynchrony
  - Complex flow of control, mutual recursion, and atomic operations
  - Dependence of a successor on its predecessor for lock passing
  - Reentry of an aborted thread into the protocol
- Result
  - T&S lock: 3 LOC
  - MCS lock: 11 LOC
  - Abortable HMCS lock: 300+ LOC



# Challenges Developing the Abortable HMCS Lock

- Ensuring the correctness of the protocol
  - Algorithm and implementation
- What happened?
  - A live lock in a simple test
- Nature of the concurrency bug:
  - Non-deterministic
  - Sporadic
  - Heisenbug
  - Many threads needed to reproduce (typically)



# Failed Efforts to Pinpoint the Bug

First instinct: “My code must be wrong!”

- Low level techniques
  - Debugging: single stepping and watch points
  - Extensive set of assertions
  - Page protection to catch stray accesses
- High level strategies
  - Protocol/design review with synchronization experts
  - Reasoning about the memory model and instruction reordering

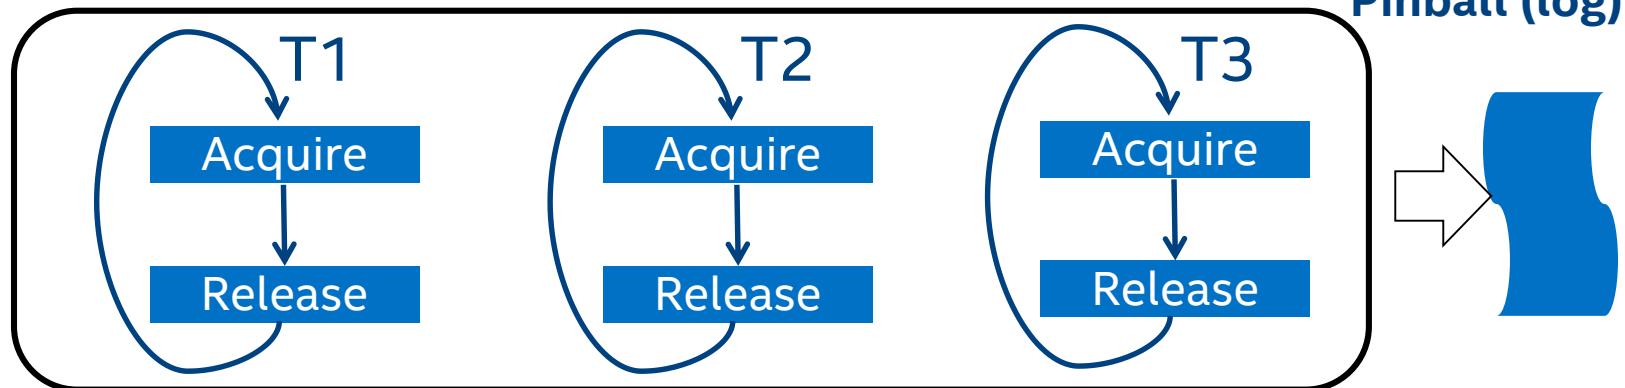
# PinPlay to Pinpoint the Bug



Record the buggy execution with PinPlay

Test setup

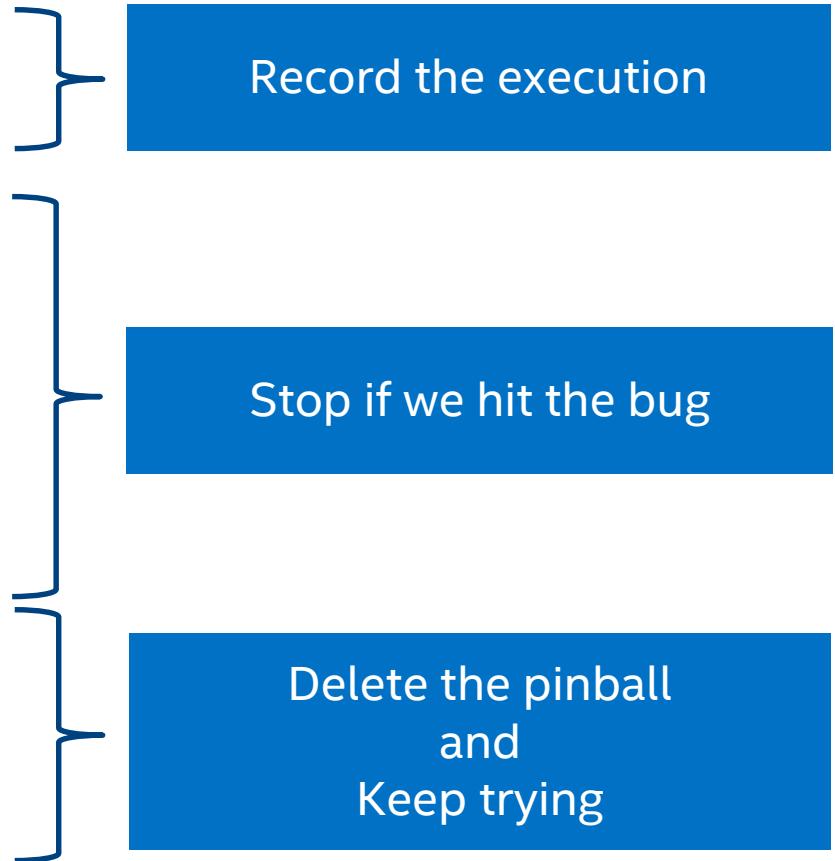
R  
e  
c  
o  
r  
d



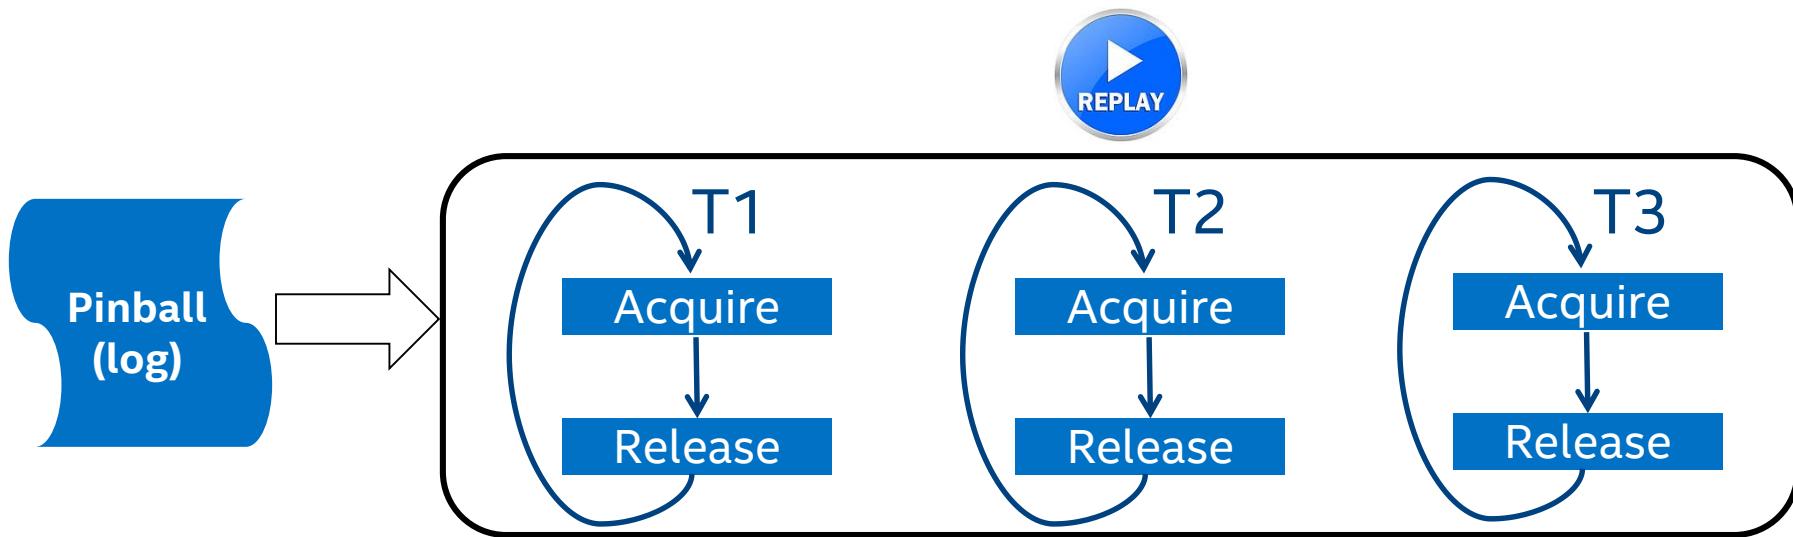
- Launch the locking test under PinPlay recorder
- If the assertion is hit then done
- Else, delete the pinball and try again
- Pinball with buggy execution ready overnight

# PinPlay Makes Execution Recoding Really Simple

```
i=0
while record ./buggy_app
do
    if test -e core*
    then
        echo "Core dumped"
        exit 1
    fi
    i=`expr $i + 1`
    echo "==== $i ==="
    rm -rf pinball
done
```

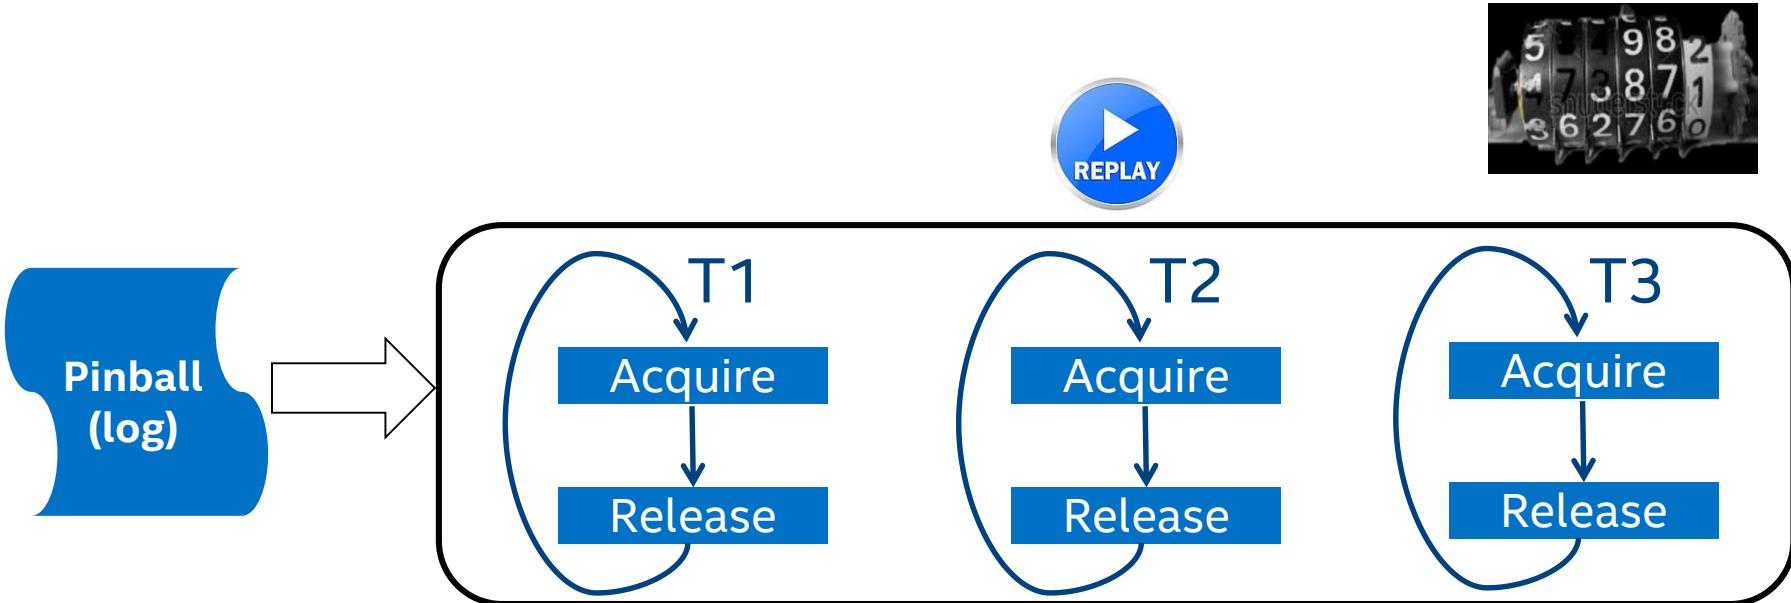


# Deterministic Debugging with Replay



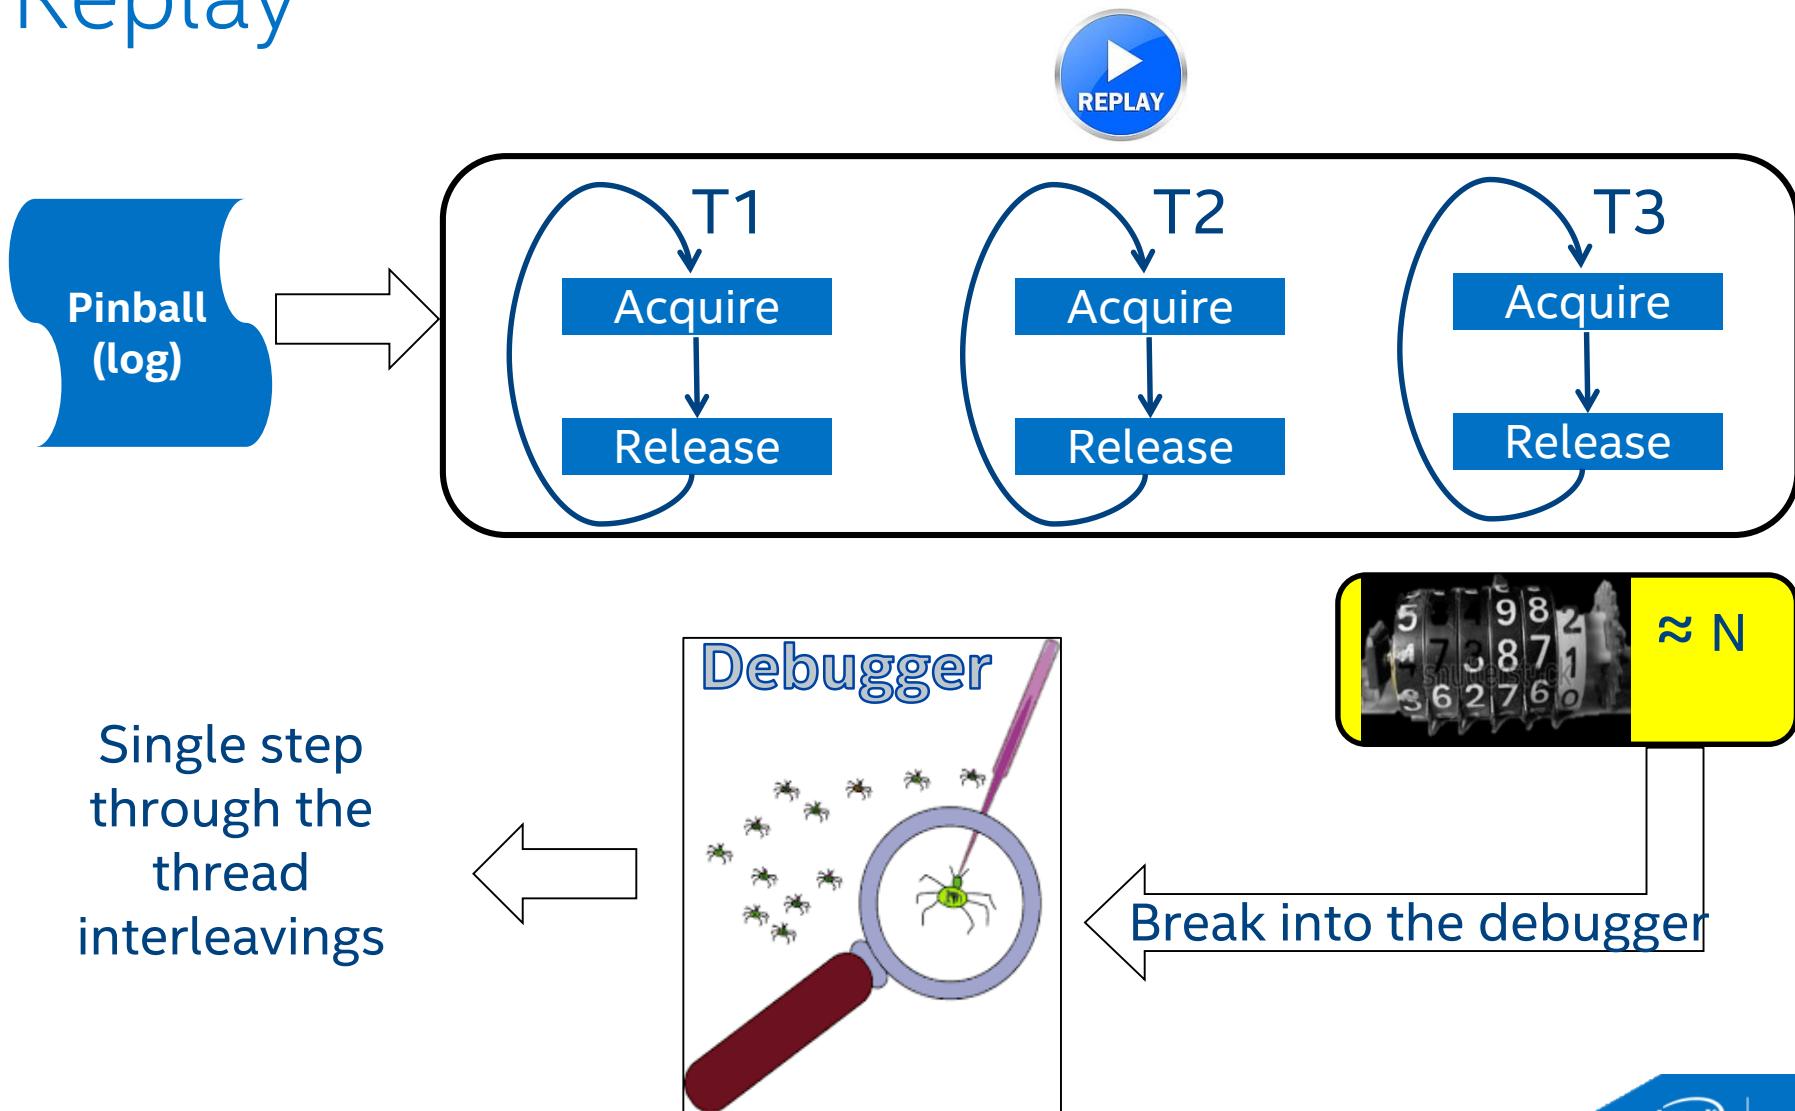
- **Challenge:** in which loop iteration did the bug happen?
  - Each thread had executed > 500K iterations
  - GDB conditional break points: too slow for remote debugging

# Deterministic Debugging with Replay

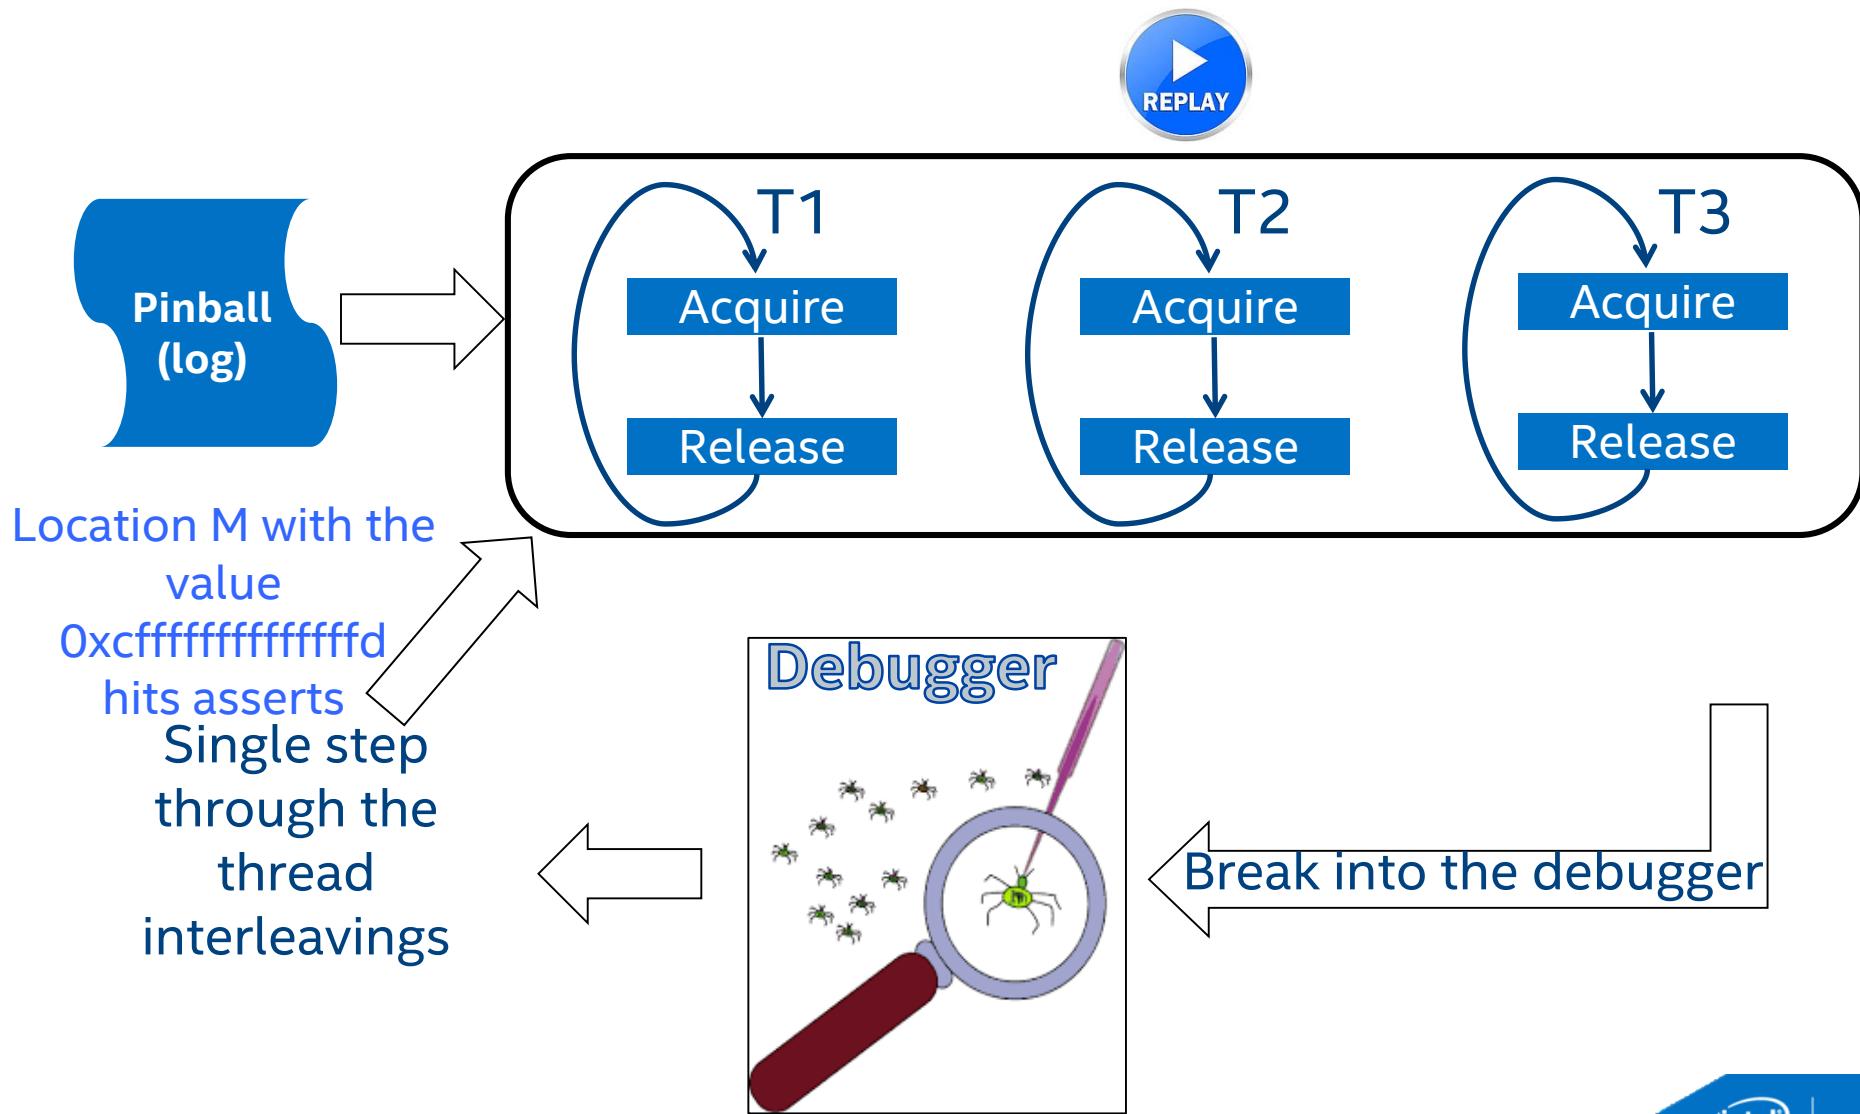


- **Challenge:** in which loop iteration did the bug happen?
  - Each thread had executed > 500K iterations
  - GDB conditional break points: too slow for remote debugging
- **Solution:** Pin tools to count/intercept/monitor interesting events

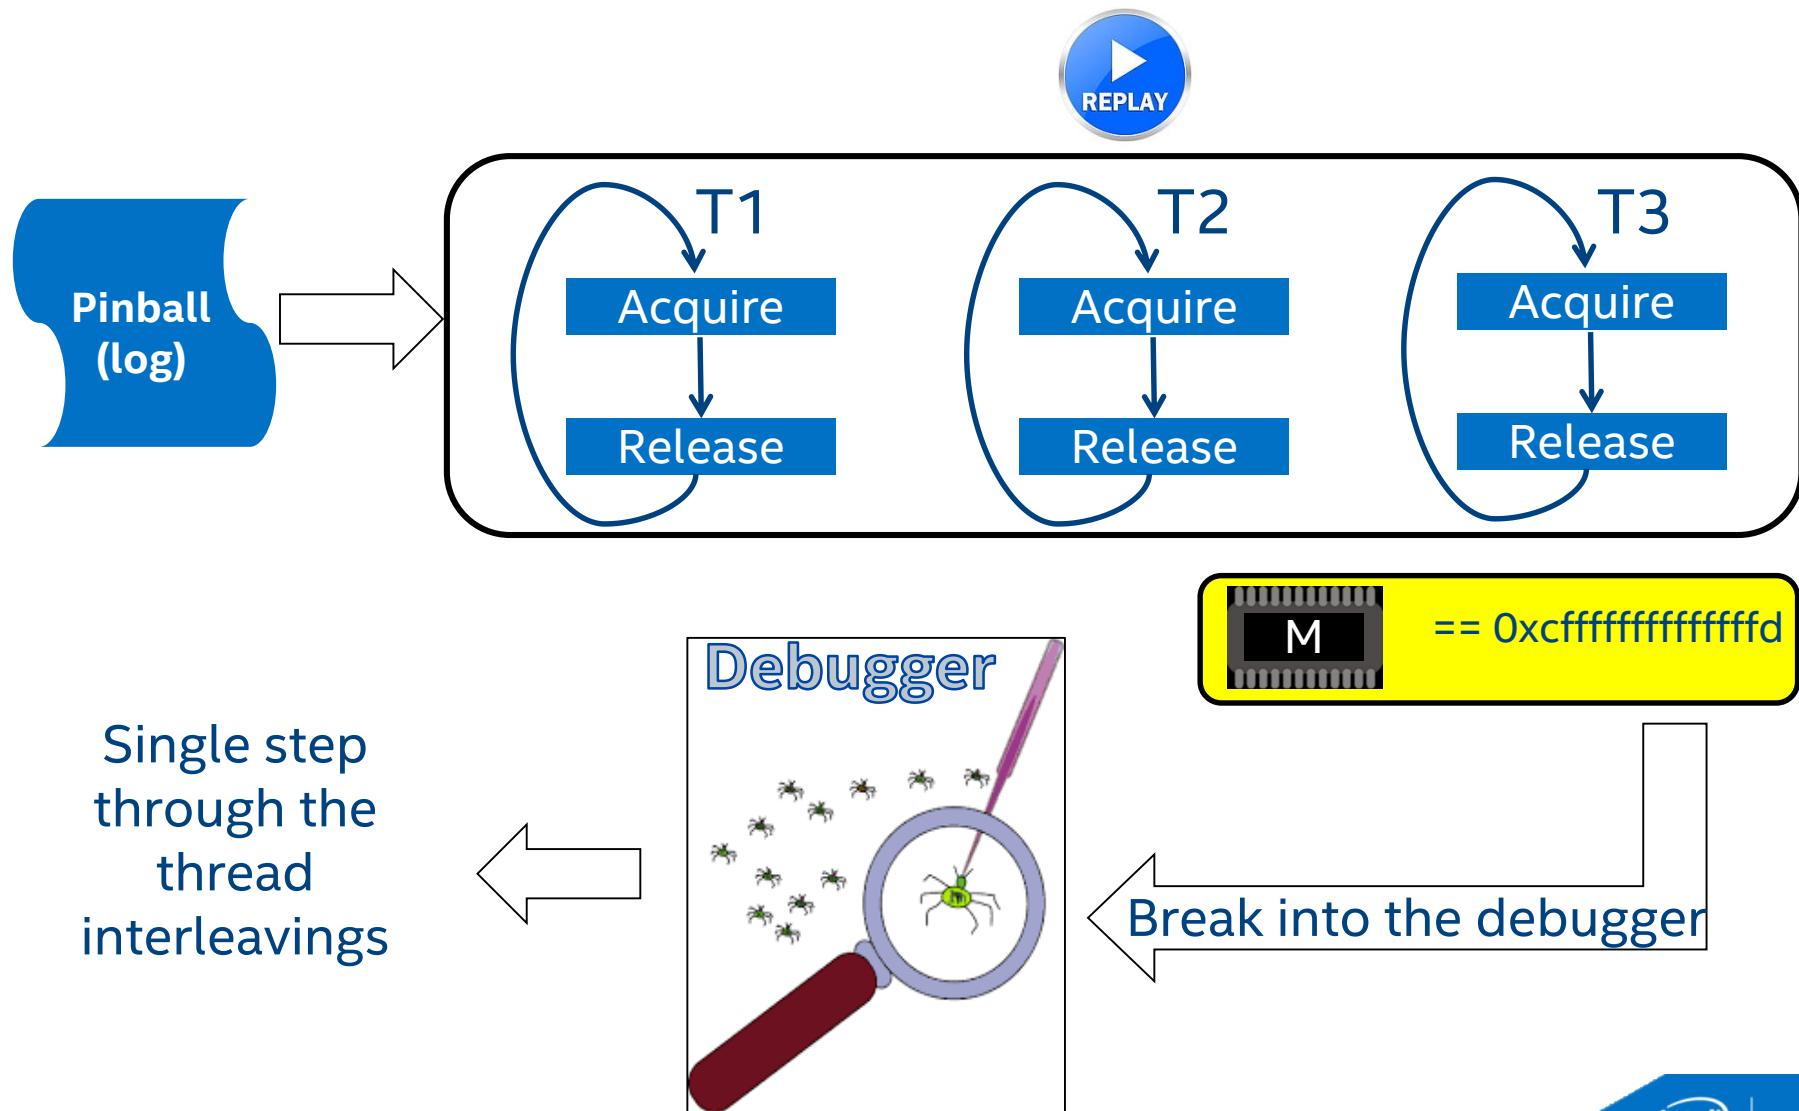
# Access to Rich Set of Pin APIs During Replay



# Effective Debugging with Faithful Replay



# Breaking into the Debugger Just In Time



# Dubious g++ Code Generation Corrupts Algorithm!

## ■ Source code

```
#define WAIT          (0xffffffffffffffff)
#define READY         (0xdfffffff00000000)
#define ACQUIRE        (0xcfffffff0000000c)
L->flag = READY // flag is 64-bit and cache line aligned
```

T1

Start with L->flag: 0xcfffffff0000000c

T2

```
movl $0xffffffff, (%rax)
//L->flag: 0xcfffffff0000000d
```

```
movl $0xdfffffff, 0x4(%rax)
//L->flag: 0xdfffffff0000000f
```

```
val = SWAP(&L->flag, WAIT)
//L->flag: 0xffffffffffff0000

switch(val) {
    case READY: /* do something */ break;
    ...
    default:
        while (L->flag != READY);
        // T1 will never make L->flag READY
}
```

# PinPlay was the Savior

- **Problem:** a non-deterministic concurrency bug
  - Not noticeable at the source (the point of write) or sync (read a clobbered, yet valid, 64-bit value)
  - Not reproducible with single-stepping the source
- **Reason:** a code generation defect that split a 64-bit write into two 32-bit writes, creating a small window of inconsistent state
- **Takeaway:** a reliable record/replay tool is invaluable in identifying concurrency bugs

# Dynamic Program Slicing with Replay



Joint work with **Yan Wang, Rajiv Gupta, and Julian Neamtiu**  
University of California, Riverside  
Eclipse GUI by **David Wootton (Intel Corporation)** sponsored by  
DARPA

# Program Slicing

Definition:  $\text{Slice}(v@S)$

Slice of  $v$  at  $S$  is the set of statements involved in computing  $v$ 's value at  $S$ . [Mark Weiser, 1982]

**Static slice** is the set of statements that **COULD** influence the value of a variable for **ANY** input

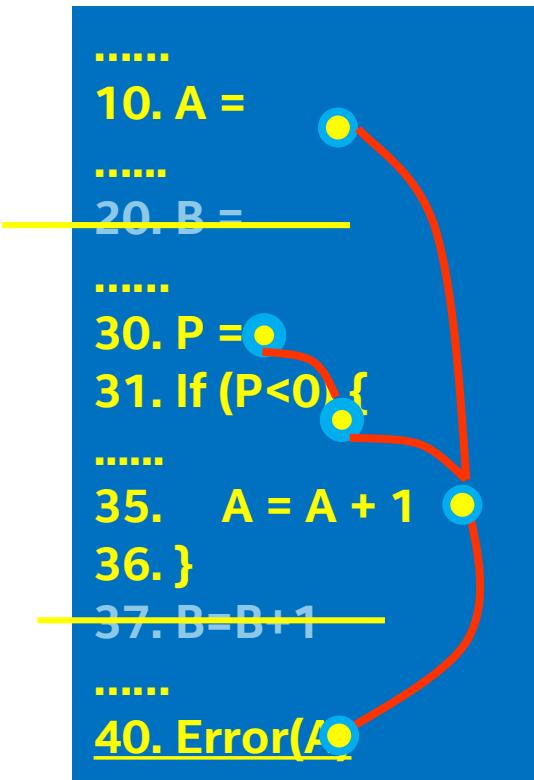
- Construct **static** dependence graph
  - Control dependences
  - Data dependences
- Traverse dependence graph to compute slice
  - Transitive closure over control and data dependences

Problem : Too many statements in a slice

# Dynamic Slicing

**Dynamic slice** is the set of statements that **DID** affect the value of a variable at a program point for **ONE specific** execution.

[Korel and Laski, 1988]



**Dynamic Slice (A@40) = {10, 30, 31, 35, 40}**

# Dynamic Slicing: Requirements

## 1. Collect Execution trace

control flow trace -- dynamic control dependences

memory reference trace -- dynamic data dependences

## 2. Construct a **dynamic** dependence graph

## 3. Traverse dynamic dependence graph to compute slices

**Challenges** : Run-time and Space overhead

# Goals of Slicing Research

*Practical Dynamic Slicing for Replay-based Interactive Debugging*

*Focus on a recording of a ‘buggy region’:*

- ❑ Manageable overhead
- ❑ Repeatability : Slice once → use across multiple debug sessions
- Develop a set of commands that integrate dynamic slicing with interactive debugging

# Distribution model : PinPlay kit + Eclipse GUI

Eclipse GUI

[www.pinplay.org](http://www.pinplay.org) → Dynamic Slicing

**PinPlay kit**

@ [www.pinplay.org](http://www.pinplay.org)

+

libpinplay.a

+

pinplay-driver  
+ debugger-shell  
+ Python scripts

+

libslicing.a

# Slicing with PinPlay: Steps

Only  
once

**Static  
Analysis:**  
computes  
STATIC CFG  
+ branch  
targets

Only  
once

**Record  
a  
region**

Only  
once

**Replay  
+  
Slicing**

As  
needed

**Replay  
+  
Debug  
+  
View  
Slice**

# Time Overhead for PARSEC

## Slicing time overhead:

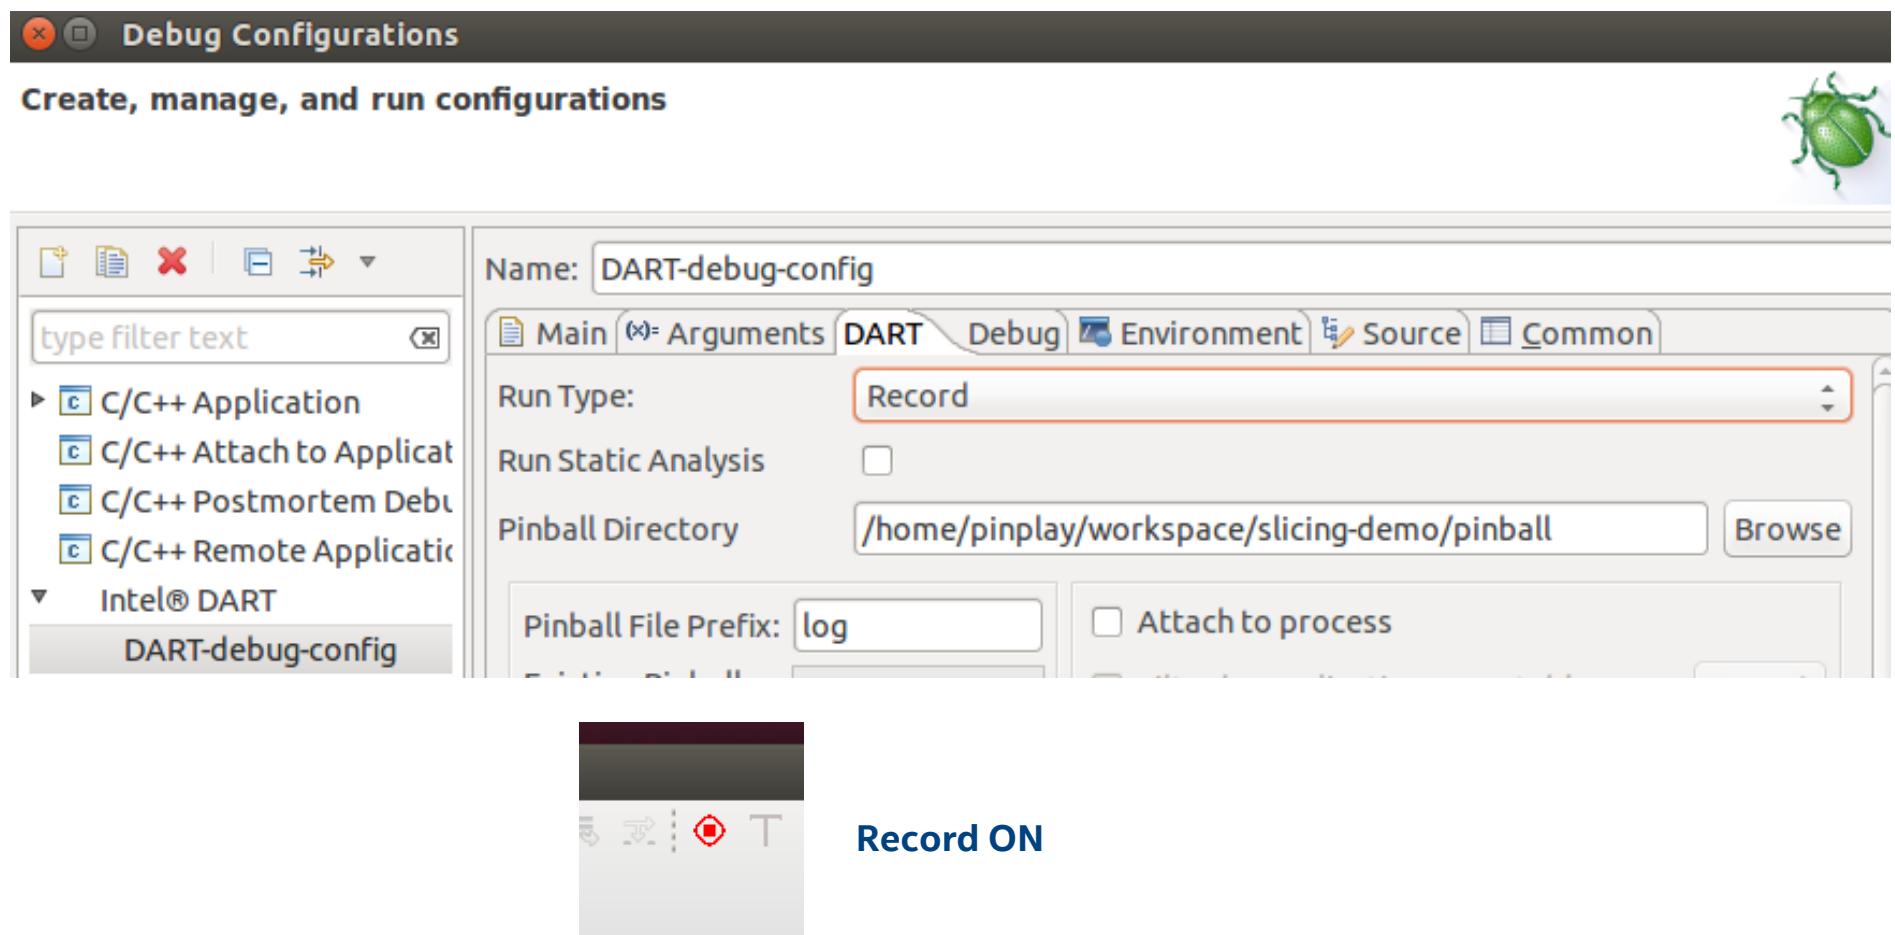
- 10 slices for the last 10 different read instructions, spread across five threads, for region length **1 million instructions (main thread)**
- Average dynamic information tracing time: **51** seconds
- Average size of slice: **218K** dynamic instructions
- Average slicing time: **585** seconds

# Record with eclipse + PinPlay

```
pinplay-VirtualBox:~/DrDebug/Slicing> EclipseDART/eclipse
```

Screenshot of the Eclipse IDE showing the "Debug Configurations" dialog. The configuration is named "DART-debug-config". The "Run Type" is set to "Record". The "Pinball Directory" is set to "/home/pinplay/workspace/slicing-demo/pinball". The "Pinball File Prefix" is set to "log". The "Attach to process" checkbox is unchecked.

The status bar at the bottom shows a green bug icon and the text "Record ON".



# Replay + Slicing with eclipse + PinPlay

Debug Configurations  
Create, manage, and run configurations



Name: DART-debug-config

Main Arguments DART Debug Environment Source Common

Run Type: Replay

Allow Slice Analysis

Run Static Analysis

Pinball Directory /home/pinplay/workspace/slicing-demo/pinball

Pinball File Prefix: log

Existing Pinballs log\_0  Refresh

Slice Criteria

Slices

|               |         |             |     |
|---------------|---------|-------------|-----|
| Instance      | 1       | Thread ID   | 1   |
| File Name     |         | Line Number | 143 |
| Start Address |         | End Address |     |
| Variable      | revenue | Length      |     |

Generate Remove Prune Show Code

intel

# Replay + Slicing with eclipse + PinPlay (contd.)

Slice Criteria   Dependency Summary   **Dependency View** ✖

Slice data for instance 1 thread 1 at line bread-demo.cpp:143 for variable revenue

```
▶ bread-demo.cpp:143 thread 1 instance 1
▶ bread-demo.cpp:141 thread 1 instance 2
▶ bread-demo.cpp:135 thread 1 instance 2
▶ bread-demo.cpp:212 thread 2 instance 14
```

135 **for (w = 0; w < NumWorkers; w++)**
136 {
137 **pthread\_join(workers[w].\_tid, 0);**
138
139 **for (i = 0; i < NUMBER\_BREADS; i++)**
140 **breadOrders[i] += workers[w].\_breadOrders[i];**
141 **revenue += workers[w].\_totalRevenue;**
142 }

# Schedule

~~9:00 – 9:45 PinPlay and PinADX~~

~~9:45 – 10:15 DrDebug demo~~

~~10:15 – 10:30 Testimonial ( Milind )~~

~~10:30 – 11:00 Slicing demo~~

<11 – 11:20 Break>

11:20 – 11:50 DCFG (Chuck)

11:50 – 12:15 Maple (Cristiano)

12:15 – 12:30 Wrap-up + Q&A (all)

# Dynamic Control-Flow Graph (DCFG)



**Chuck Yount**  
**Intel Corporation**

# Overview

## Tutorial goals

- Show how to create a Dynamic Control-Flow Graph (DCFG) from a PinPlay-enabled tool
- Illustrate additional control-flow analysis enabled by PinPlay's ability to replay execution

## Agenda

- DCFG Definition
- How to make a DCFG
- How to use a DCFG
- Optional DCFG-Trace creation and usage

DCFG web site for documentation and download

- <https://software.intel.com/en-us/articles/pintool-dcfg>
- Or, <http://pinplay.org> and follow the DCFG link

# DCFG definition

## Control-Flow Graph [Allen 1970] (CFG)

- Directed graph in which nodes represent basic blocks and edges represent control-flow paths
- Basic block: linear sequence of instructions having one entry point and one exit point
- Used in many compiler and analysis algorithms

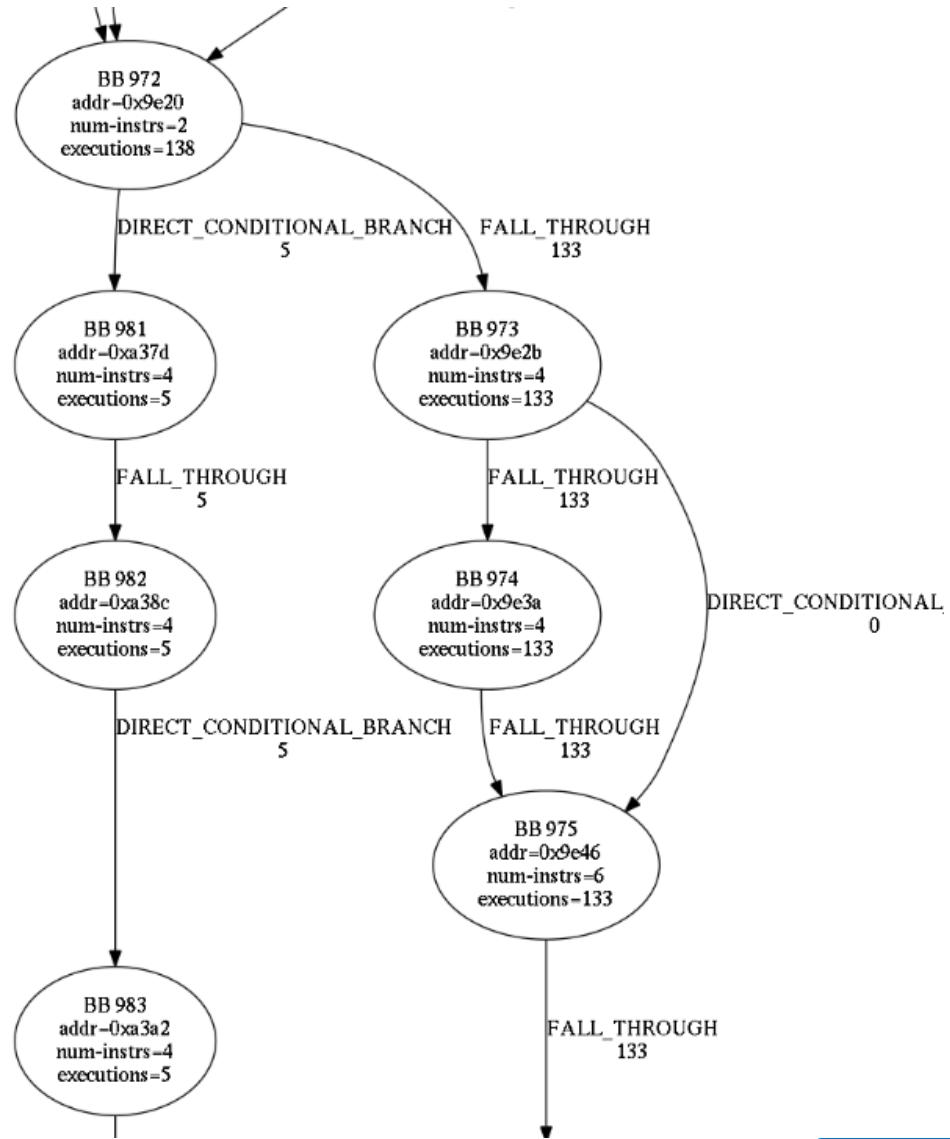
## Dynamic Control-Flow Graph (DCFG)

- Defined by and extracted from a particular execution of a program
- Edges augmented with per-thread execution count; basic-block and other counts can be derived from these
- Need not contain non-executed edges or blocks
- May include paths due to exceptions, etc.

# Example DCFG snippet

## Shows a nested conditional

- BB (basic block) 972 entered 138 times
  - If-then-else construct
  - Conditional branch to BB 981 (left side) taken 5 times
  - Fall-through to BB 973 remaining 133 times
    - If-then construct
    - Fall-through to BB 974 always taken
- This image was created using the 'dcfg-to-dot' utility program included in the package



# Prerequisite tutorial setup

## Virtual machine

- Install Oracle VM VirtualBox or equivalent VM
- Download the PinPlay virtual appliance and import into the VM

## Or, local Linux install

- Download the PinPlay package from <http://pinplay.org>
- Extract the tarball
- Set PIN\_ROOT and PIN\_HOME environment vars to the install directory

# How to create a DCFG

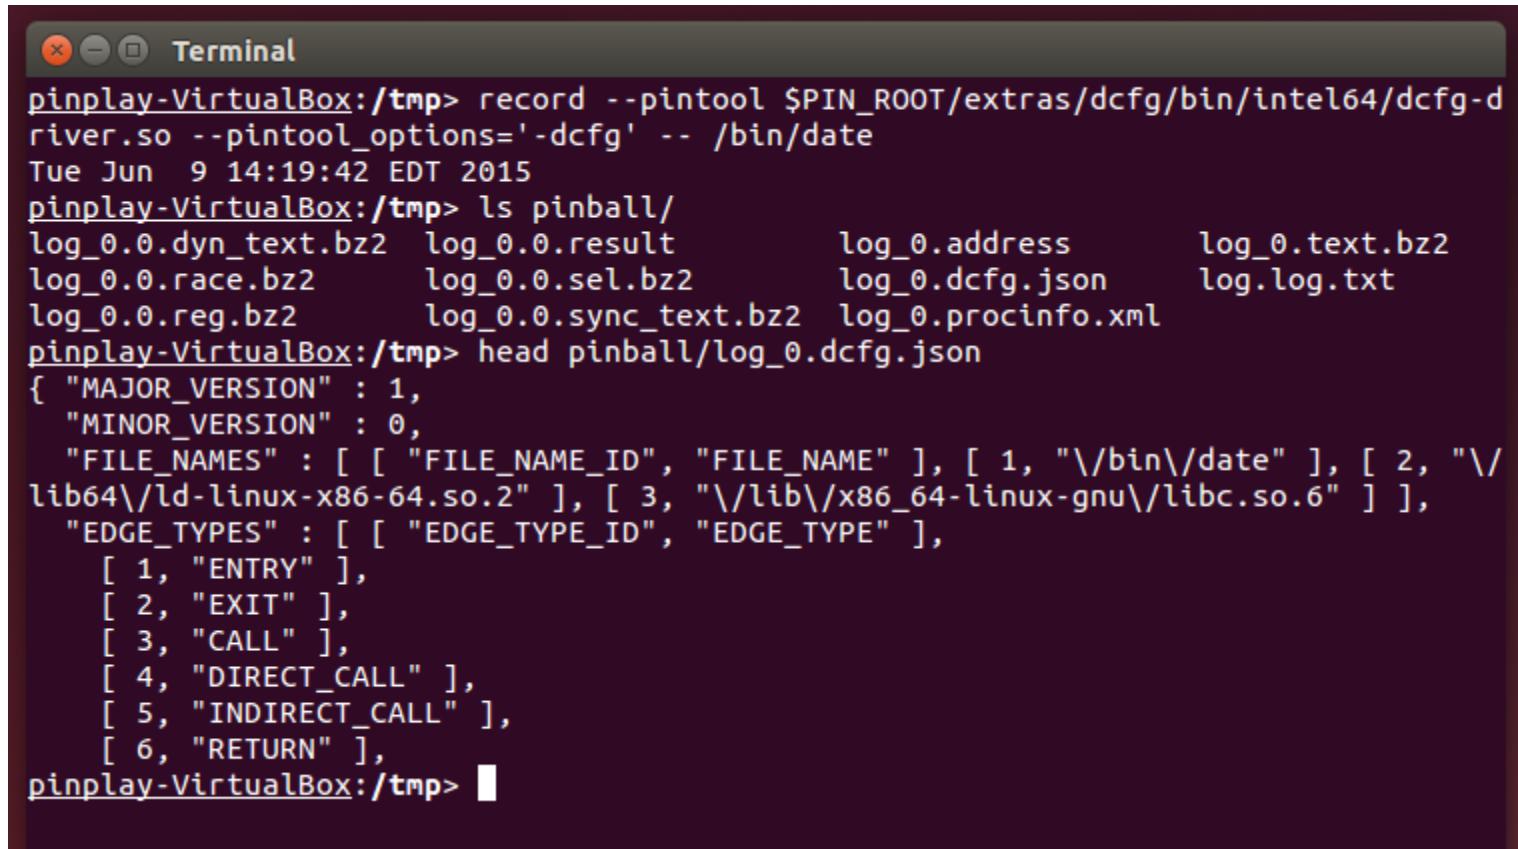
Run the ‘record’ script with the DCFG driver tool

- `record --pintool  
$PIN_ROOT/extras/dcfg/bin/intel64/dcfg-driver.so --  
pintool_options='"-dcfg'" -- /bin/date`

Creates new file in ‘pinball’ directory: `log_0.dcfg.json`

- Contains data on logged process in JSON format
  - Meta-data: Images, symbols, and debug info
  - Fundamental CFG elements: basic blocks and edges
  - Derived structures: routines and loops
  - See “DCFG format description” on website for full documentation

# DCFG creation example



The screenshot shows a terminal window titled "Terminal" with the following command history:

```
pinplay-VirtualBox:/tmp> record --pintool $PIN_ROOT/extras/dcfg/bin/intel64/dcfdriver.so --pintool_options='--dcfg' -- /bin/date
Tue Jun  9 14:19:42 EDT 2015
pinplay-VirtualBox:/tmp> ls pinball/
log_0.0.dyn_text.bz2  log_0.0.result          log_0.address      log_0.text.bz2
log_0.0.race.bz2       log_0.0.sel.bz2        log_0.dcfg.json    log.log.txt
log_0.0.reg.bz2        log_0.0.sync_text.bz2  log_0.procinfo.xml
pinplay-VirtualBox:/tmp> head pinball/log_0.dcfg.json
{
  "MAJOR_VERSION" : 1,
  "MINOR_VERSION" : 0,
  "FILE_NAMES" : [
    [ "FILE_NAME_ID", "FILE_NAME" ],
    [ 1, "\/bin\/date" ],
    [ 2, "\lib64\ld-linux-x86-64.so.2" ],
    [ 3, "\lib\x86_64-linux-gnu\libc.so.6" ]
  ],
  "EDGE_TYPES" : [
    [ "EDGE_TYPE_ID", "EDGE_TYPE" ],
    [ 1, "ENTRY" ],
    [ 2, "EXIT" ],
    [ 3, "CALL" ],
    [ 4, "DIRECT_CALL" ],
    [ 5, "INDIRECT_CALL" ],
    [ 6, "RETURN" ]
  ]
}
pinplay-VirtualBox:/tmp>
```

# DCFG-creation code

In \$PIN\_ROOT/extras/dcfg

- See examples/makefile.rules
- examples/dcfg-driver.cpp provides minimal DCFG functionality

```
#include "dcfg_pin_api.H" ...  
DCFG_PIN_MANAGER* dcfgMgr =  
DCFG_PIN_MANAGER::new_manager();  
if (dcfgMgr->dcfg_enable_knob())  
    dcfgMgr->activate(&pinplay_engine);
```

- Link with lib/arch/libdcfg-pinplay.a
  - Provides '-dcfg' and other DCFG command-line options

# Reading a DCFG file (standalone tool)

C++ API usable from standalone program or from a PinPlay tool

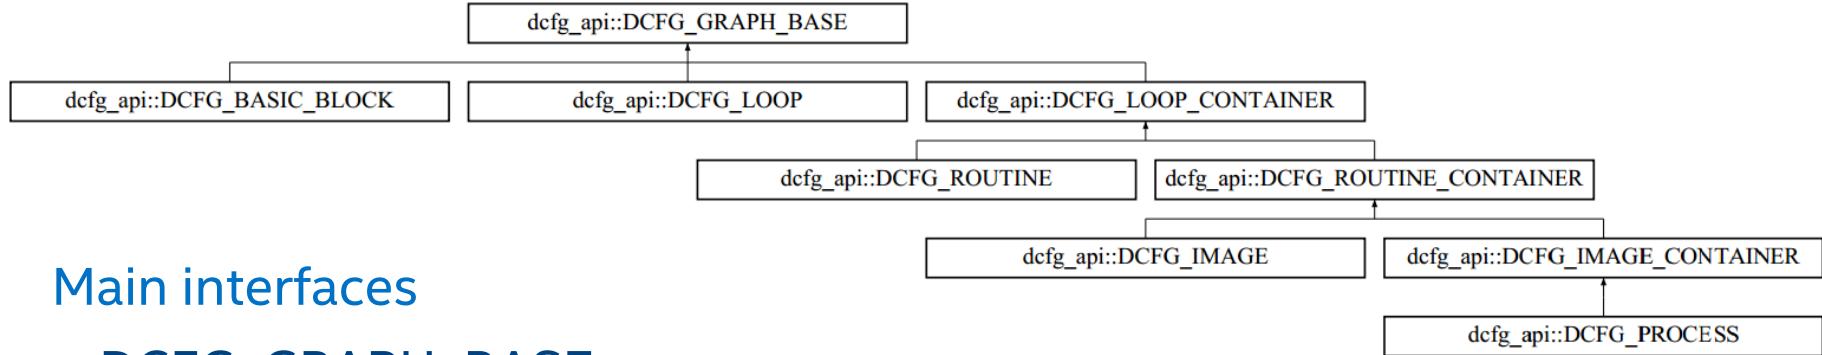
- Documentation at DCFG web site
  - “Hierarchical Index” is a good starting point
- Example standalone code in examples/dcfg-reader.cpp
  - Link with lib/arch/libdcfg.a
  - Create a DCFG\_DATA object and read contents from a file

```
#include "dcfg_api.H "
DCFG_DATA* dcfg = DCFG_DATA::new_dcfg();
dcfg->read(filename, errMsg);
```
  - Most data is accessed by getting one or more IDs, e.g.,  
`dcfg->get_process_ids(proc_ids);`
  - Then, use an ID to get a pointer to detailed data, e.g.,  
`DCFG_PROCESS_CPTR pinfo =
dcfg->get_process_info(proc_ids[i]);`
  - Similar code to get data on images, routines, loops, basic blocks, edges, etc.
- Run example code to print high-level statistics
  - `$PIN_ROOT/extras/dcfg/bin/intel64/dcfg-reader pinball/log_0.dcfg.json`

# Standalone-code example

```
Terminal
pinplay-VirtualBox:/tmp> $PIN_ROOT/extras/dcfg/bin/intel64/dcfg-reader pinball/log_0.dcfg.json
| head -25
Reading DCFG from 'pinball/log_0.dcfg.json'...
Summary of DCFG:
Num processes          = 1
Process 2266
  Num threads = 1
  Instr count = 222435
  Num edges   = 4238
  Num images  = 3
Image 1
  Load addr      = 0x400000
  Size           = 2155712
  File           = '/bin/date'
  Num basic blocks = 220
  Num routines    = 41
  Num loops       = 4
Image 2
  Load addr      = 0x7ff1f0dff000
  Size           = 2245064
  File           = '/lib64/ld-linux-x86-64.so.2'
  Num basic blocks = 1278
  Num routines    = 74
  Num loops       = 82
Image 3
  Load addr      = 0x7ff1d8c07000
  Size           = 3949248
  File           = '/lib/x86_64-linux-gnu/libc.so.6'
```

# API interface inheritance hierarchy



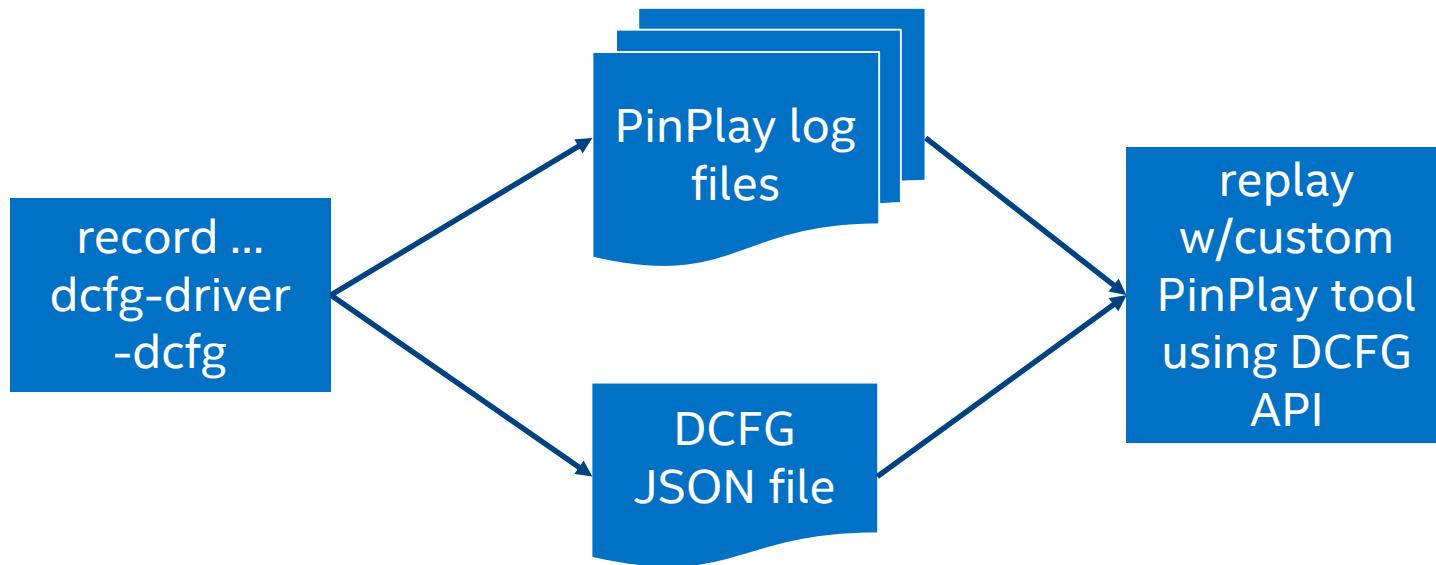
## Main interfaces

- **DCFG\_GRAPH\_BASE**
  - Base interface for anything with basic blocks and edges
- **DCFG\_{LOOP,ROUTINE,IMAGE}\_CONTAINER**
  - Base interface for something containing the given part
  - Example: a process and an image are both routine containers
- **DCFG\_{BASIC\_BLOCK,LOOP,ROUTINE,IMAGE,PROCESS}**
  - Main hierarchical structural components
- **DCFG\_EDGE (not shown)**
  - Control-flow path between any two nodes (usually BBs)

# Using a DCFG during replay

Provides access to DCFG structures during replay

- Can instrument code based on analysis of edges, loops, etc.

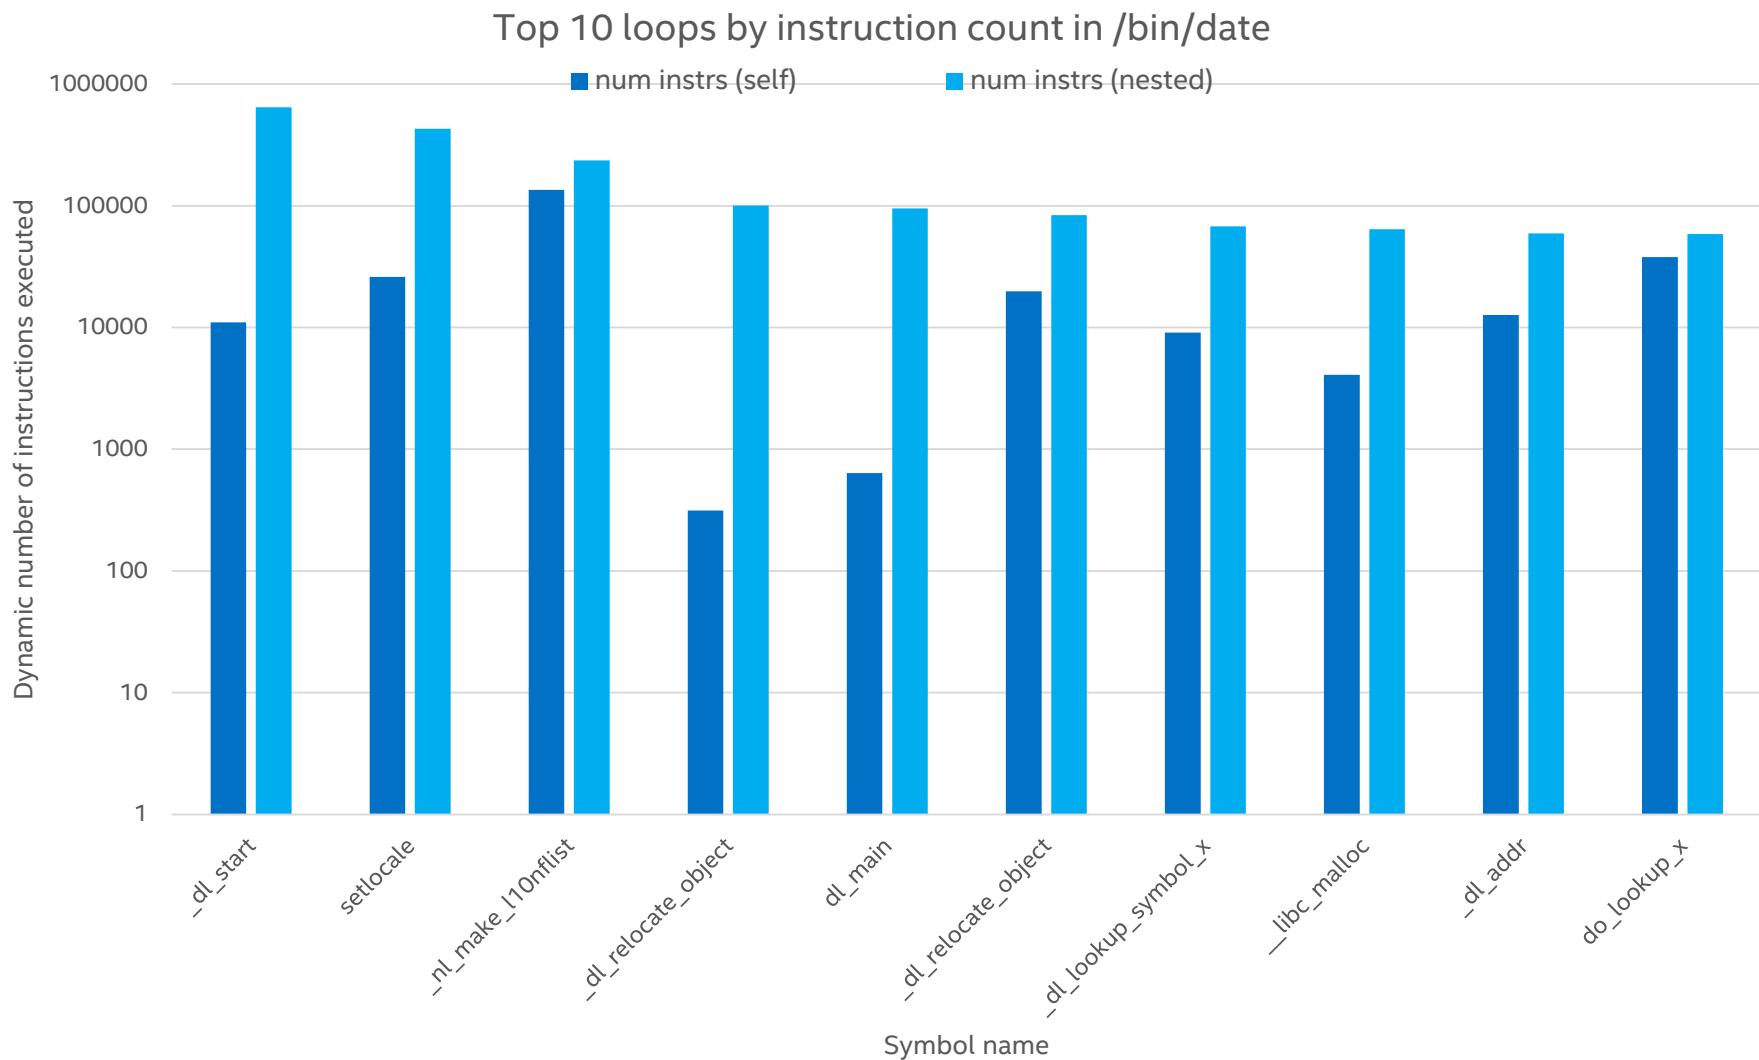


# Reading a DCFG file (PinPlay tool)

Same C++ API can be used in a PinPlay tool

- Example code in examples/loop-tracker.{cpp,H}
  - Collects statistics from loops
  - Link with lib/arch/libdcfg-pinplay.a as before
  - Noteworthy methods in LOOP\_TRACKER class
    - activate(): reads the DCFG file and adds instrumentation functions
    - processDcfg(): collects data from the DCFG data to be used later
    - handleTrace(): adds analysis routine at DCFG basic block heads
    - enterBb(): analysis code that is run at each basic block
      - Remembers previous BB to determine DCFG edge
      - Uses data from processDcfg() to track loop entry and exit
      - Collects statistics on each loop
    - printData(): outputs statistics for each loop in tabular format
- Run loop tracker
  - ```
replay --pintool
$PIN_ROOT/extras/dcfd/bin/intel64/loop-tracker.so --
pintool_options='-loop-tracker:dcfg-file
pinball/log_0.dcfg.json' -- pinball/log_0
```
  - Creates loop-stats.csv by default

# Data from example loop-tracker run



# DCFG PinPlay tool example

```
Terminal
pinplay-VirtualBox:/tmp> replay --pintool $PIN_ROOT/extras/dcfg/bin/intel64/loop-tracker.so --p
intoool_options=' -loop-tracker:dcfg-file pinball/log_0.dcfg.json' -- pinball/log_0
Tue Jun  9 14:19:42 EDT 2015
pinplay-VirtualBox:/tmp> head loop-stats.csv
loop id,thread id,image name,symbol name,source file,source line number,loop addr,num loop entr
ies,num iterations,ave iterations/entry,ave loop-nesting depth,num instrs (self),num instrs (ne
sted),ave num instrs/iteration (self),ave num instrs/iteration (nested)
0,0,"/lib64/ld-linux-x86-64.so.2",".text","unknown",0,0x7ff1f0e03a70,1,1,1.00,0.00,9379,222467,
9379.00,222467.00
100,0,"/bin/date","**entry**","unknown",0,0x402e79,1,15,15.00,1.00,3658,6464,243.87,430.93
144,0,"/bin/date","**entry**","unknown",0,0x4034c3,5,11,2.20,2.00,182,182,16.55,16.55
180,0,"/bin/date","**entry**","unknown",0,0x4041d0,1,1,1.00,2.00,86,86,86.00,86.00
218,0,"/bin/date","**entry**","unknown",0,0x409390,1,1,1.00,1.00,22,22,22.00,22.00
231,0,"/lib64/ld-linux-x86-64.so.2",".text","unknown",0,0x7ff1f0ffff45,1,4,4.00,1.00,24,24,6.00
,6.00
274,0,"/lib64/ld-linux-x86-64.so.2",".text","unknown",0,0x7ff1f0e01239,1,9,9.00,1.00,156,156,17
.33,17.33
296,0,"/lib64/ld-linux-x86-64.so.2",".text","unknown",0,0x7ff1f0e01409,1,23,23.00,1.00,201,201,
8.74,8.74
336,0,"/lib64/ld-linux-x86-64.so.2",".text","unknown",0,0x7ff1f0e02032,1,1,1.00,1.00,8,8,8.00,8
.00
```

# Utility to create visual DCFG

Sample tool to output the DCFG of one routine in 'DOT' format

- Usage information
  - `$PIN_ROOT/extras/dcfg/bin/intel64/dcfg-to-dot`  
Writes graph of routine containing node-id to a DOT-language file.  
usage: `dcfg-to-dot <input-dcfg-file> <output-dot-file> <node-id>`
- The 'node-id' can be any node in the routine for which you want to create the dot-file
- Use 'dot' program to convert to PDF, PNG, etc.
  - Download and install from <http://www.graphviz.org/>

# DCFG-Trace

Additional optional record of entire sequence of edges

- Enabled by using -dcfg:write\_trace option
- Allows standalone-tool to analyze sequence
  - Use DCFG\_TRACE\_READER API
  - See summarizeTrace() in examples/dcfg-reader.cpp
    - Inputs DCFG-Trace file
    - Outputs table of edges and basic-blocks at edge targets
- PinPlay tool provides many more possibilities (all Pin APIs) and does not require disk-space for large edge-trace file

# DCFG-Trace example

```
Terminal
pinplay-VirtualBox:/tmp> record --pintool $PIN_ROOT/extras/dcfg/bin/intel64/dcfg-driver.so --pintool_options='-dcfg -dcfg:write_trace' -- /bin/date
Tue Jun  9 20:13:35 EDT 2015
pinplay-VirtualBox:/tmp> $PIN_ROOT/extras/dcfg/bin/intel64/dcfg-reader pinball/log_0.dcfg.json
pinball/log_0.trace.json | grep -A10 'edge id'
Reading DCFG from 'pinball/log_0.dcfg.json'...
Reading DCFG trace for PID 2921 and TID 0 from 'pinball/log_0.trace.json'...
edge id,basic-block id,basic-block addr,basic-block symbol,num instrs in BB
2549,410,0x7fbe0bf6cabf,".text",1
2727,411,0x7fbe0bf6cac5,".text",11
2285,414,0x7fbe0bf6cb1f,".text",2
307,412,0x7fbe0bf6cb08,".text",2
309,413,0x7fbe0bf6cb13,".text",4
300,414,0x7fbe0bf6cb1f,".text",2
307,412,0x7fbe0bf6cb08,".text",2
309,413,0x7fbe0bf6cb13,".text",4
300,414,0x7fbe0bf6cb1f,".text",2
637,415,0x7fbe0bf6cb25,".text",4
Done reading 56096 edges.
```

# Summary

## DCFG creation

- Minimal dcfg-driver.so PinPlay tool creates DCFG JSON file
  - Contains structure of basic-blocks, edges, loops, and more
- Optionally, DCFG-Trace file can be created
  - Contains sequence of edges

## DCFG usage

- Example dcfg-reader standalone program inputs DCFG (and optionally DCFG-Trace) and prints summary data
- Example loop-tracker.so PinPlay tool inputs DCFG and instruments basic-blocks to track edges and loops

# Triggering Bugs for Repeatable Analysis



**Cristiano Pereira**  
Intel Corporation

# Overview

- Tutorial goals
  - Describe a tool to trigger an existing concurrency error in a multi-threaded program execution
  - Demonstrate how PinPlay can be used to capture an execution that contains an exposed concurrency error
- Agenda
  - Maple: A tool and technique to trigger concurrency errors in a multi-threaded program
  - Brief overview of how to use Maple
  - How to capture a bug exposed by Maple with PinPlay and replay-debug it

# Triggering Bugs for Repeatable Analysis

- Replay is a powerful technique for analyzing the root cause of concurrency bugs
- However, the bug must manifest during recording for replay to be applicable, and this may not be easy
- A mechanism for triggering concurrency bugs is desirable
  - Single-threaded: mostly driven by program input
  - Multi-threaded programs: driven by both program input and by thread-interleaving
  - Thread-interleaving depends on # of cores, cache sizes, memory latencies, OS scheduler, machine load, position of stars in the sky...

# Motivation

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <assert.h>

unsigned NUM_THREADS = 1;
unsigned global_count = 0;
void *thread(void *);

int main(int argc, char *argv[]) {
    pthread_t *pthread_id[100];
    NUM_THREADS = atoi(argv[1]);

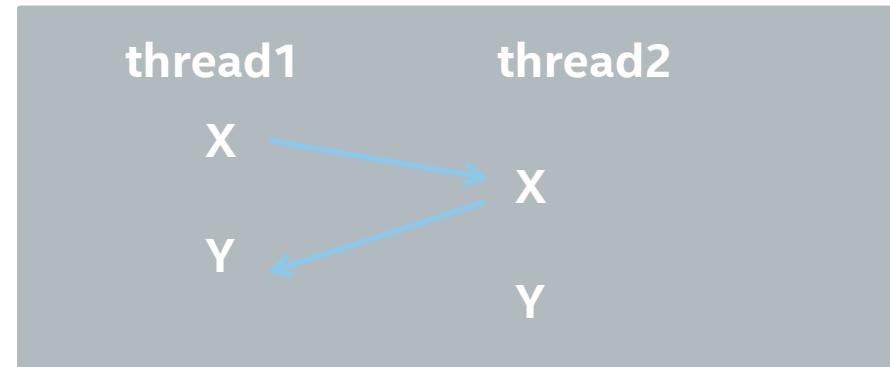
    for(long i = 0; i < NUM_THREADS; i++)
        pthread_create(&pthread_id[i],
                      NULL, thread, (void *) i);

    for(long i = 0; i < NUM_THREADS; i++)
        pthread_join(pthread_id[i], NULL);

    assert(global_count==NUM_THREADS);
    exit(0);
}
```

```
void * thread(void * num)
{
    X: unsigned temp = global_count;
        temp++;
    Y: global_count = temp;
        return NULL;
}
```

What is the buggy interleaving?



How frequently?

# crashes/10,000 runs (8 cores):

0 with 8 threads/run (0%)

2 with 32 threads/run (0.02%)

# Maple Tool

- Goal:
  - Expose executions that exhibit buggy behavior
- How:
  - Finds memory operations likely to lead to bugs
  - Produces a buggy run by manipulating thread schedule using Linux real time priorities
  - Implemented in PIN, easy to integrate with other pintools

# Triggering Concurrency Bugs with Maple

Maple: a coverage-driven testing tool for multithreaded programs, Yu et al. OOPSLA'12  
<https://github.com/jieyu/maple>

- Maple uses a “best effort” approach by profiling the execution before hand and looking for specific bug patterns or idioms
  - An idiom is a sequence of memory operations and their dependencies across threads
  - Limits the search by observing that bugs are not caused by random memory operations
  - No completeness guarantees like systematic (formal) testing, but a practical technique that scales to relatively large programs
  - Details in the OOPSLA'12 paper, out of scope for this tutorial

# Idioms of bugs with 2 threads, $\leq 2$ variables

*Idiom-1*



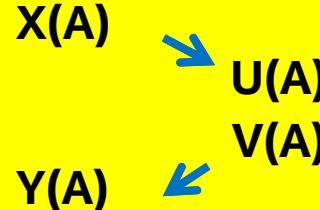
Data race or ordering

*Idiom-2*



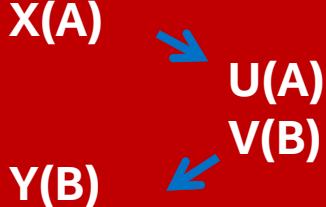
Low level atomicity

*Idiom-3*



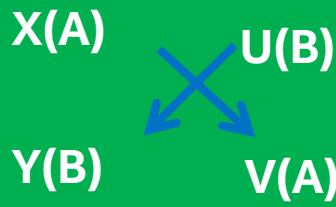
Single variable

*Idiom-4*



Multi-variable atomicity

*Idiom-5*



Deadlock

*Idiom-6*



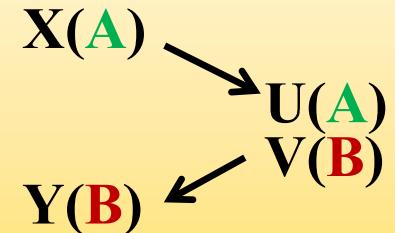
Double variable

Unknown

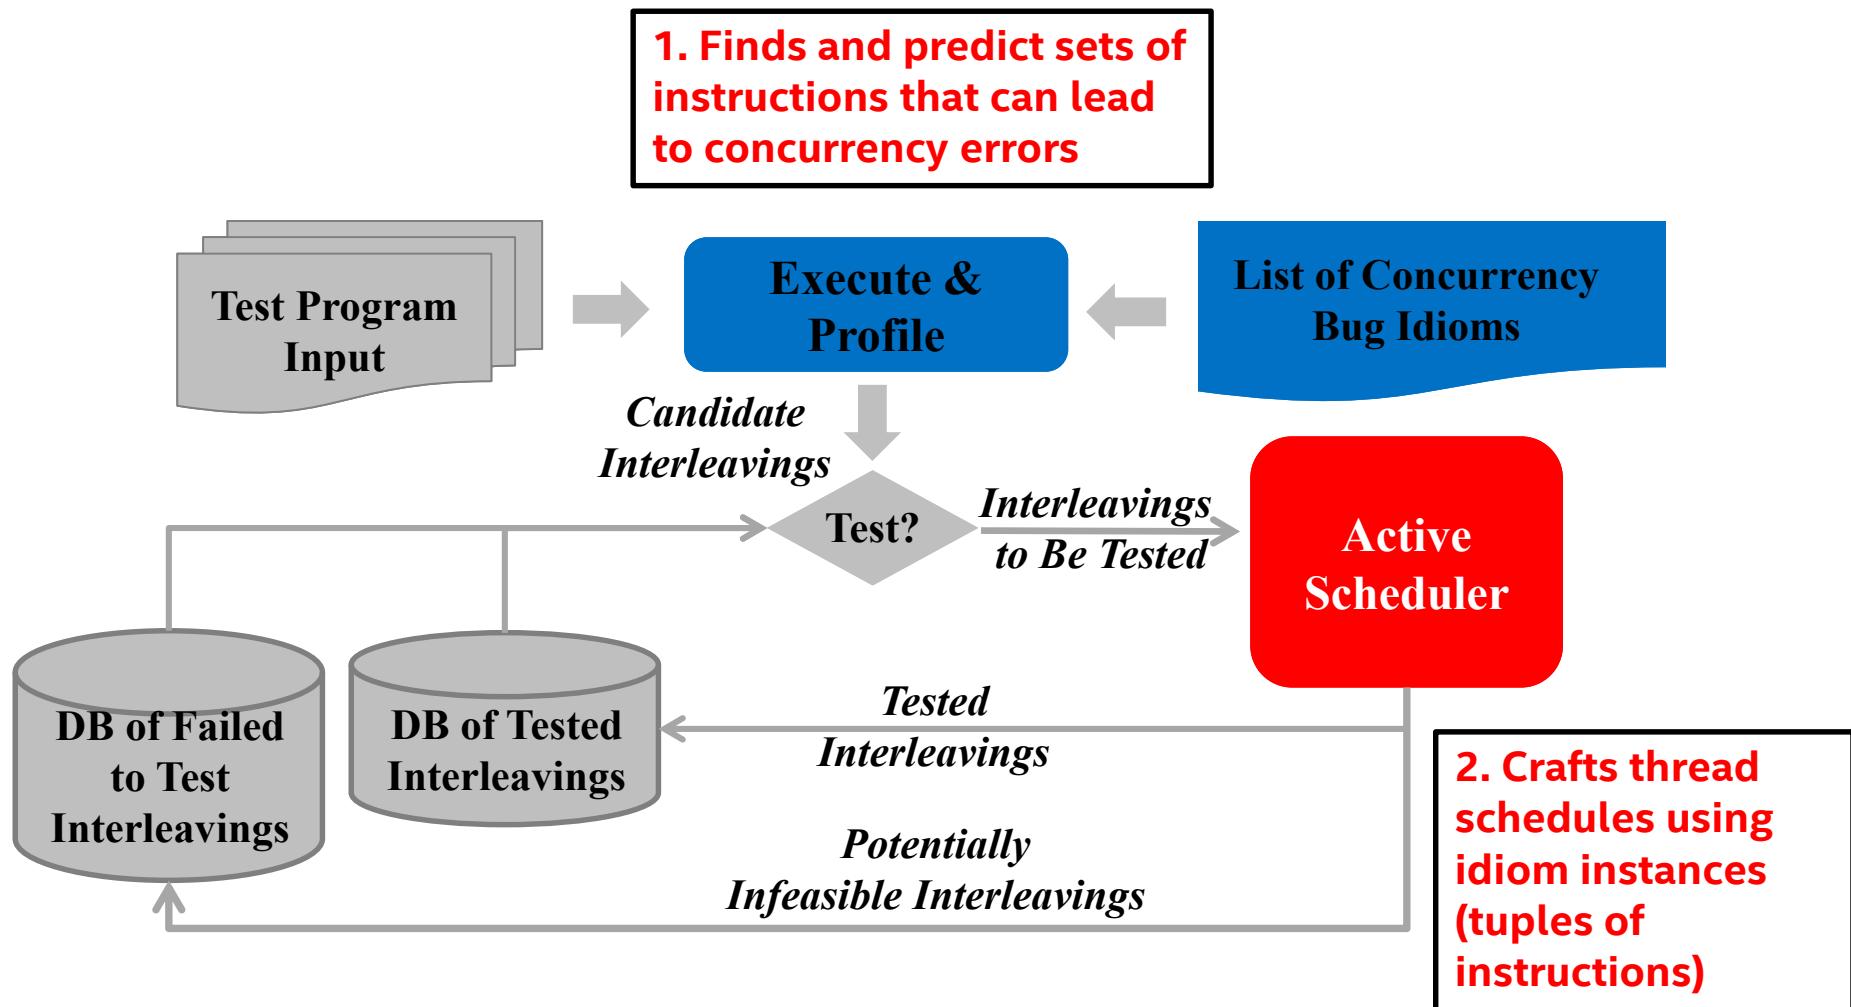
# Idiom Example: MySQL Bug #169

```
int generate_table(...)  
{  
    lock(&LOCK_open);  
    // delete entries  
X: unlock(&LOCK_open);  
  
Y: lock(&LOCK_log);  
    log.write(...);  
    unlock(&LOCK_log);  
  
    V: lock(&LOCK_log);  
    log.write(...);  
    unlock(&LOCK_log);  
  
    U: lock($LOCK_open);  
    // insert entry  
    unlock($LOCK_open);  
}  
int mysql_insert(...)  
{  
}
```

**Template:**



# Maple Bug Triggering Overview



# Exposing Idiom Instances

- An *idiom* instance, called an *iroot*, consists of the PCs of memory operations that fit within the idiom

```
2      IDIOM_1
      e0: WRITE   [79    ld      0x4c9f
      e1: READ    [1212   ld      0x9f7e ]
```

- A data base of *iroots* is built during profiling
- Maple then executes the program multiple times forcing each *iroot* to be exercised during the scheduling phase
- If an execution exposes bug, it is reproduced by repeating the *iroot* inter-thread dependencies

# Maple and PinPlay Usage

- Maple exercises all interleavings in the iroot database
- When maple reproduces a thread interleaving, it only reproduces the interleaving involved in the idiom that exposed the bug
- All other sources of non-determinism are not controlled
  - Memory addresses may differ
  - Control-flow may differ
  - Interleaving is not guaranteed to repeat 100% of the time
- Combining Maple and PinPlay guarantees that an execution exposing the idiom instance can be repeated 100% deterministically and replay-debugged that way

# Modified Shared Counter Example

Change lines as highlighted in the example program under:

/home/pinplay/maple/example/shared\_counter/main.cc

```
17 #include <stdio.h>
18 #include <stdlib.h>
19 #include <pthread.h>
20 #include <assert.h>
21
22 unsigned LOOP_COUNT = 100;
23 unsigned NUM_THREADS = 1;
24 unsigned global_count = 0;
25 void *thread(void *);
26
27 int main(int argc, char *argv[]) {
28     long i;
29     pthread_t pthread_id[200];
30     NUM_THREADS = atoi(argv[1]);
31
32     for(i = 0; i < NUM_THREADS; i++)
33         pthread_create(&pthread_id[i], NULL, thread, (void *) i);
34     for(i = 0; i < NUM_THREADS; i++)
35         pthread_join(pthread_id[i], NULL);
36
37     printf("global_count: %d\n", global_count);
38     assert(global_count==(NUM_THREADS*LOOP_COUNT));
39     return 0;
40 }
41
42 void *thread(void * num) {
43     for (int i=0; i < LOOP_COUNT; i++)
44         global_count++;
45 }
46
```

# Running Maple

```
Terminal
pinplay-VirtualBox:~/maple/example/shared_counter> ../../maple --profile_single_var_idioms --- ./main
2
global_count: 200
[MAPLE] === profile iteration 1 done === (1.412212) (/home/pinplay/maple/example/shared_counter)
global_count: 200
[MAPLE] === profile iteration 2 done === (1.609778) (/home/pinplay/maple/example/shared_counter)
global_count: 200
[MAPLE] === profile iteration 3 done === (1.409103) (/home/pinplay/maple/example/shared_counter)
global_count: 200
[MAPLE] === profile iteration 4 done === (1.509463) (/home/pinplay/maple/example/shared_counter)
global_count: 200
[MAPLE] === profile iteration 5 done === (1.713700) (/home/pinplay/maple/example/shared_counter)
[MAPLE] profile threshold reached
global_count: 100
main: main.cc:38: int main(int, char**): Assertion `global_count==(NUM_THREADS*LOOP_COUNT)' failed.
[MAPLE] === active iteration 1 done === (1.109534) (/home/pinplay/maple/example/shared_counter)
[MAPLE] active fatal error detected
[MAPLE]
[MAPLE] -----
[MAPLE] profile_runs      5
[MAPLE] profile_time      7.854002
[MAPLE] active_runs       1
[MAPLE] active_time       1.125862
pinplay-VirtualBox:~/maple/example/shared_counter>
```

Profiling  
and  
building  
iroots from  
predefined  
idioms

Actively  
scheduling  
threads to  
reproduced  
the iroots  
found

# Visualizing the IRoot database

```
pinplay-VirtualBox:~/maple/example/shared_counter> ../../script/idiom display iroot_db
```

|    |         |           |       |      |         |             |   |
|----|---------|-----------|-------|------|---------|-------------|---|
| 1  | IDIOM_1 | e0: WRITE | [1206 | ld   | 0xa8cf  | ]           |   |
|    |         | e1: READ  | [1206 | ld   | 0xa8cf  | ]           |   |
| 2  | IDIOM_1 | e0: WRITE | [79   | ld   | 0x4c9f  | ]           |   |
|    |         | e1: READ  | [1212 | ld   | 0x9f7e  | ]           |   |
| 3  | IDIOM_1 | e0: WRITE | [1365 | ld   | 0x100c1 | ]           |   |
|    |         | e1: READ  | [1214 | ld   | 0x9fa4  | ]           |   |
| 4  | IDIOM_1 | e0: WRITE | [70   | ld   | 0xb972  | ]           |   |
|    |         | e1: READ  | [1216 | ld   | 0x9fbe  | ]           |   |
| 5  | IDIOM_1 | e0: WRITE | [8    | ld   | 0x4b0f  | ]           |   |
|    |         | e1: READ  | [1217 | ld   | 0x9fcc  | ]           |   |
| 6  | IDIOM_1 | e0: WRITE | [8    | ld   | 0x4b0f  | ]           |   |
|    |         | e1: READ  | [1219 | ld   | 0x9fd4  | ]           |   |
| 7  | IDIOM_1 | e0: WRITE | [75   | ld   | 0xb995  | ]           |   |
|    |         | e1: READ  | [1221 | ld   | 0x9fe1  | ]           |   |
| 8  | IDIOM_1 | e0: WRITE | [73   | ld   | 0xb986  | ]           |   |
|    |         | e1: READ  | [1222 | ld   | 0x9ff4  | ]           |   |
| 9  | IDIOM_1 | e0: WRITE | [76   | ld   | 0xb99c  | ]           |   |
|    |         | e1: READ  | [1224 | ld   | 0x9ffe  | ]           |   |
| 10 | IDIOM_1 | e0: WRITE | [77   | ld   | 0xb9ad  | ]           |   |
|    |         | e1: READ  | [1226 | ld   | 0xa794  | ]           |   |
| 28 | IDIOM_2 | e0: WRITE | [1206 | ld   | 0xa8cf  | ]           |   |
|    |         | e1: WRITE | [1206 | ld   | 0xa8cf  | ]           |   |
|    |         | e2: READ  | [1206 | ld   | 0xa8cf  | ]           |   |
| 29 | IDIOM_1 | e0: READ  | [1754 | main | 0x859   | main.cc +44 | ] |
|    |         | e1: WRITE | [1755 | main | 0x862   | main.cc +44 | ] |
| 30 | IDIOM_1 | e0: WRITE | [1755 | main | 0x862   | main.cc +44 | ] |
|    |         | e1: WRITE | [1755 | main | 0x862   | main.cc +44 | ] |

# Reproducing the Bug with Maple

```
pinplay-VirtualBox:~/maple/example/shared_counter> ../../script/idiom display test_history  
29 IDIOM_1 Success 1433996132
```

Last iroot tested, ID=29  
and random seed=1433996132

```
pinplay-VirtualBox:~/maple/example/shared_counter> ../../script/idiom active --target_iroot=29  
--random_seed=1433996132 --- ./main 2  
global_count: 100  
main: main.cc:38: int main(int, char**): Assertion `global_count==(NUM_THREADS*LOOP_COUNT)' failed.  
[MAPLE] === active iteration 1 done === (1.113106) (/home/pinplay/maple/example/shared_counter)  
[MAPLE] active fatal error detected
```

Bug reproduced by  
Maple by repeating same  
interleaving

# Capturing the Bug With PinPlay

```
pinplay-VirtualBox:~/maple/example/shared_counter> ../../script/idiom active --target_iroot=29 --random_seed=1433996132 --pinplay --pinplay_options="-log:basename,failing.pinball/log" --- ./main 2
global_count: 100
main: main.cc:38: int main(int, char**): Assertion `global_count==(NUM_THREADS*LOOP_COUNT)' failed.
[MAPLE] === active iteration 1 done === (6.938914) (/home/pinplay/maple/example/shared_counter)
[MAPLE] active fatal error detected
```

Execution captured in a log file

```
pinplay-VirtualBox:~/maple/example/shared_counter> replay failing.pinball/log
global_count: 100
main: main.cc:38: int main(int, char**): Assertion `global_count==(NUM_THREADS*LOOP_COUNT)' failed.
pinplay-VirtualBox:~/maple/example/shared_counter> replay failing.pinball/log
global_count: 100
main: main.cc:38: int main(int, char**): Assertion `global_count==(NUM_THREADS*LOOP_COUNT)' failed.
pinplay-VirtualBox:~/maple/example/shared_counter> replay failing.pinball/log
global_count: 100
main: main.cc:38: int main(int, char**): Assertion `global_count==(NUM_THREADS*LOOP_COUNT)' failed.
pinplay-VirtualBox:~/maple/example/shared_counter>
```

Execution replayed deterministically

# Replaying within GDB

```
pinplay-VirtualBox:~/maple/example/shared_counter> gdb_replay failing.pinball/log main
GNU gdb (Ubuntu 7.7.1-0ubuntu5~14.04.2) 7.7.1
Copyright (C) 2014 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word"...
Reading symbols from main...done.
_dl_start (arg=0x0) at rtld.c:362
362      rtld.c: No such file or directory.
(gdb) b main.cc:38
Breakpoint 1 at 0x400080f: file main.cc, line 38.
(gdb) c
Continuing.
global_count: 100

Breakpoint 1, main (argc=2, argv=0x7fffffa3e188) at main.cc:38
38      assert(global_count==(NUM_THREADS*LOOP_COUNT));
(gdb) print global_count
$1 = 100
(gdb) c
Continuing.
main: main.cc:38: int main(int, char**): Assertion `global_count==(NUM_THREADS*LOOP_COUNT)' failed.
[Inferior 1 (Remote target) exited normally]
(gdb) █
```

# Using GDB + PinPlay to Debug

- PinPlay interfaces with PinADX and is enabled with a set of new GDB commands

```
(gdb) monitor help
general           - General commands.
breakpoints       - Breakpoint commands.
tracepoints       - Tracepoint commands.
registers         - Register names.
PinPlay           - PinPlay commands.

type "help <category>" for help on commands in that category.
```

- Among the new commands, PinPlay is enabled with tracepoints, allowing the capture of a trace during execution (triggered by PC value, load or store value, etc)
- In this tutorial, we will use tracepoints to inspect the value of the `global_counter` and how it is updated by various threads

# Modifying Tracepoint Implementation

- Tracepoint implementation can be found at:  
*/home/pinplay/PinPlay/pinplay-drdebug-2.2-pldi2015-pin-2.14-71313-gcc.4.4.7-linux/extras/pinplay/examples/pinplay-debugger-shell.cpp*
- We will use the trace point to capture memory addresses executed at a given PC (the read and the write to `global_count`)

```
trace <addr> <length> at <pc>
    Record trace entry at <pc> from memory value from
    'addr' of <length> bytes.
```

- However, the default implementation does not output the thread id for each trace, so we modify it to do so

```
struct TRACEREC
{
    unsigned _id;          // Index of EVENT in '_events'.
    ADDRINT _pc;           // PC where tracepoint triggered.
    ADDRINT _req_mem_value; // If tracepoints traces a register, it's value.
    THREADID _tid;         // Thread ID added for PLDI 2015 tutorial
};
```

\* Analysis functions to collect tracepoints are also changed to take the thread ID

# Creating a Tracepoint

```
(gdb) b main
Breakpoint 1 at 0x400745: file main.cc, line 30.
(gdb) c
Continuing.

Breakpoint 1, main (argc=2, argv=0x7fff3e67a2e8) at main.cc:30
30      NUM_THREADS = atoi(argv[1]);
(gdb) pin trace global_count at 44
monitor trace 0x60106c 4 at 0x400859 #global_count:<44>
Tracepoint #1: trace memory address 0x00060106c length 4
  at 0x400859 #global_count:<44>
(gdb) b 38
Breakpoint 2 at 0x40080f: file main.cc, line 38.
(gdb) c
Continuing.
global_count: 100

Breakpoint 2, main (argc=2, argv=0x7fff3e67a2e8) at main.cc:38
38      assert(global_count==(NUM_THREADS*LOOP_COUNT));
(gdb) pin trace print to gcount.txt
monitor trace print to gcount.txt
```

- File gcount.txt contains the traces collected at runtime
- Now we need to inspect it!

# Inspecting the Tracepoint File

|                                                                                                                                            |                                                                                                            |
|--------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------|
| 2 0x00000000000400859: 0x00060106c = 0x0                                                                                                   | 1 Thread 2 reads initial counter value, context switches out                                               |
| 3 0x00000000000400859: 0x00060106c = 0x0                                                                                                   |                                                                                                            |
| 3 0x00000000000400868: 0x00060106c = 0x1                                                                                                   |                                                                                                            |
| 3 0x00000000000400859: 0x00060106c = 0x1                                                                                                   |                                                                                                            |
| 3 0x00000000000400868: 0x00060106c = 0x2                                                                                                   |                                                                                                            |
| 3 0x00000000000400859: 0x00060106c = 0x2                                                                                                   |                                                                                                            |
| 3 0x00000000000400868: 0x00060106c = 0x3                                                                                                   |                                                                                                            |
| 3 0x00000000000400859: 0x00060106c = 0x3                                                                                                   |                                                                                                            |
| 3 0x00000000000400868: 0x00060106c = 0x4                                                                                                   |                                                                                                            |
| 3 0x00000000000400859: 0x00060106c = 0x4                                                                                                   |                                                                                                            |
| 3 0x00000000000400868: 0x00060106c = 0x5                                                                                                   |                                                                                                            |
| 3 0x00000000000400859: 0x00060106c = 0x5                                                                                                   |                                                                                                            |
| 3 0x00000000000400868: 0x00060106c = 0x6                                                                                                   |                                                                                                            |
| 3 0x00000000000400859: 0x00060106c = 0x6                                                                                                   |                                                                                                            |
| 3 0x00000000000400868: 0x00060106c = 0x7                                                                                                   |                                                                                                            |
| 3 0x00000000000400859: 0x00060106c = 0x7                                                                                                   |                                                                                                            |
| 3 0x00000000000400868: 0x00060106c = 0x8                                                                                                   |                                                                                                            |
| 3 0x00000000000400868: 0x00060106c = 0x61                                                                                                  |                                                                                                            |
| 3 0x00000000000400859: 0x00060106c = 0x61                                                                                                  |                                                                                                            |
| 3 0x00000000000400868: 0x00060106c = 0x62                                                                                                  |                                                                                                            |
| 3 0x00000000000400859: 0x00060106c = 0x62                                                                                                  |                                                                                                            |
| 3 0x00000000000400868: 0x00060106c = 0x63                                                                                                  |                                                                                                            |
| 3 0x00000000000400859: 0x00060106c = 0x63                                                                                                  |                                                                                                            |
| 3 0x00000000000400868: 0x00060106c = 0x64                                                                                                  |                                                                                                            |
| 2 0x00000000000400868: 0x00060106c = 0x1                                                                                                   | global_count incremented up to 100, thread 3 finishes                                                      |
| 2 0x00000000000400859: 0x00060106c = 0x1                                                                                                   |                                                                                                            |
| 2 0x00000000000400868: 0x00060106c = 0x2                                                                                                   |                                                                                                            |
| 2 0x00000000000400859: 0x00060106c = 0x2                                                                                                   |                                                                                                            |
| 2 0x00000000000400868: 0x00060106c = 0x3                                                                                                   |                                                                                                            |
| 2 0x00000000000400859: 0x00060106c = 0x3                                                                                                   |                                                                                                            |
| 2 0x00000000000400868: 0x00060106c = 0x62                                                                                                  | Thread 2 context switches back in, overwrites the last update by Thread 3 and then increments count to 100 |
| 2 0x00000000000400859: 0x00060106c = 0x62                                                                                                  |                                                                                                            |
| 2 0x00000000000400868: 0x00060106c = 0x63                                                                                                  |                                                                                                            |
| 2 0x00000000000400859: 0x00060106c = 0x63                                                                                                  |                                                                                                            |
| 2 0x00000000000400868: 0x00060106c = 0x64                                                                                                  |                                                                                                            |
| Interactive analysis of a particular run is a lot easier with deterministic replay because addresses and thread interleaving never changes |                                                                                                            |

# Summary

- Deterministic replay is a great tool for root causing concurrency errors
- However capturing a buggy execution may take a long time depending on the bug
- Active testing via Maple addresses the bug triggering issue and, combined with PinPlay, provides a framework to expose and analyze concurrency errors

# PinPlay: A framework for Deterministic Replay and Reproducible Analysis

*PinPlay is an **easy-to-use**, **flexible**, and **effective** framework for reproducible analysis of multi-threaded programs.*

- **easy-to-use** : the Pin advantage (download and use, no re-compile...)
- **Flexible**: combine with any Pintool, combine with debuggers
- **Effective**: it works (deterministic replay → reproducible analysis)

# Legal Disclaimer & Optimization Notice

INFORMATION IN THIS DOCUMENT IS PROVIDED "AS IS". NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO THIS INFORMATION INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products.

Copyright © 2014, Intel Corporation. All rights reserved. Intel, Pentium, Xeon, Xeon Phi, Core, VTune, Cilk, and the Intel logo are trademarks of Intel Corporation in the U.S. and other countries.

## Optimization Notice

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice revision #20110804

